

#CLASSIFICATION : TRUE/FALSE :

Membres: Hadjoudja Bachir (21811363), Zeggar Rym (21909615), Bendahmane Rania (21811387), Labiad Youcef (21710780).

#les imports utilisés dans ce notebook

```
import sys
from numpy import vstack
import pandas as pd
from pandas import read_csv
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from torch.utils.data import random_split
from torch import Tensor
from torch.nn import Linear
from torch.nn import ReLU
from torch.nn import Sigmoid
from torch.nn import Module
from torch.optim import SGD
from torch.nn import BCELoss
from torch.nn.init import kaiming_uniform_
from torch.nn.init import xavier_uniform_
import re
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from pandas import read_csv
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
import pickle
import string

import nltk
from nltk.stem import WordNetLemmatizer
from nltk.stem import PorterStemmer
from nltk.corpus import stopwords
from nltk import word_tokenize
from sklearn.pipeline import Pipeline

import sklearn
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
import seaborn as sns
```

```

import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_fscore_support as score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

```

```

from sklearn.metrics._plot.confusion_matrix import
ConfusionMatrixDisplay
# fonction qui affiche le classification report et la matrice de
confusion
from sklearn import metrics
from sklearn.metrics import confusion_matrix , ConfusionMatrixDisplay
from sklearn.metrics import classification_report

```

```

import warnings
warnings.filterwarnings('ignore')

```

```

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

chemin spécifique Google Drive

```

my_local_drive='/content/gdrive/My Drive/'
# Ajout du path pour les librairies, fonctions et données
sys.path.append(my_local_drive)
# Se positionner sur le répertoire associé
%cd $my_local_drive
%ls

```

```
%pwd
```

```

[Errno 2] No such file or directory: '/content/gdrive/My Drive/'
/content
drive/  sample_data/

```

```
{"type": "string"}
```

La fonction qui sera utilisée pour les prétraitements: MyCleanText

- Mettre le texte en minuscule
- Se débarrasser des stopwords
- Se débarrasser des nombres
- Stemmatisation
- Lemmatisation ..

La fonction MyshowAllScores prend le y_test et le y_predict, affiche l'accuracy et le classification report avec la matrice de confusion.

```
#.....Fonction
MyCleanText .....
.....
# mettre en minuscule
#enlever les stopwords
#se debarrasser des nombres
#stemmatisation
#lemmatisation
#.....
.....
.....

nltk.download('wordnet')
nltk.download('stopwords')
nltk.download('punkt')
#liste des stopwords en anglais
stop_words = set(stopwords.words('english'))

def MyCleanText(X,
                 lowercase=False, #mettre en minuscule
                 removestopwords=False, #supprimer les stopwords
                 removedigit=False, #supprimer les nombres
                 getstemmer=False, #conserver la racine des termes
                 getlemmatisation=False #lemmatisation des termes
                 ):
    #conversion du texte d'entrée en chaîne de caractères
    sentence=str(X)
    #suppression des caractères spéciaux
    sentence = re.sub(r'^\w\s',' ', sentence)
    # suppression de tous les caractères uniques
    sentence = re.sub(r'\s+[a-zA-Z]\s+', ' ', sentence)
    # substitution des espaces multiples par un seul espace
    sentence = re.sub(r'\s+', ' ', sentence, flags=re.I)

    # decoupage en mots
    tokens = word_tokenize(sentence)
    if lowercase:
        tokens = [token.lower() for token in tokens]

    # suppression ponctuation
    table = str.maketrans('', '', string.punctuation)
    words = [token.translate(table) for token in tokens]

    # suppression des tokens non alphanumérique ou numérique
    words = [word for word in words if word.isalnum()]
```

```

# suppression des tokens numerique
if removedigit:
    words = [word for word in words if not word.isdigit()]

# suppression des stopwords
if removestopwords:
    words = [word for word in words if not word in stop_words]

# lemmatisation
if getlemmatisation:
    lemmatizer=WordNetLemmatizer()
    words = [lemmatizer.lemmatize(word) for word in words]

# racinisation
if getstemmer:
    ps = PorterStemmer()
    words=[ps.stem(word) for word in words]

sentence= ' '.join(words)

return sentence

```

```

def MyshowAllScores(y_test,y_pred):
    classes= np.unique(y_test)
    print("Accuracy : %0.3f"%(accuracy_score(y_test,y_pred)))
    print("Classification Report")
    print(classification_report(y_test,y_pred,digits=5))
    cnf_matrix = confusion_matrix(y_test,y_pred)
    disp=ConfusionMatrixDisplay(cnf_matrix,display_labels=classes)
    disp.plot()

```

```

[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!

```

- La classe TextNormalizer qui contiendra la fonction MyCleanText.
- Fit_transform de mon corpus propre.

```

#.....Etape 1 :
prétraitement du
texte .....
.....
#.....Class

```

```

TextNormalizer .....
.....
#fit_transform de mon corpus propre
#.....
.....
.....

```

```

from sklearn.base import BaseEstimator, TransformerMixin

```

```

class TextNormalizer(BaseEstimator, TransformerMixin):
    def __init__(self,
        removestopwords=False, # suppression des stopwords
        lowercase=False, # passage en minuscule
        removedigit=False, # supprimer les nombres
        getstemmer=False, # racinisation des termes
        getlemmatisation=False # lemmatisation des termes
    ):

        self.lowercase=lowercase
        self.getstemmer=getstemmer
        self.removestopwords=removestopwords
        self.getlemmatisation=getlemmatisation
        self.removedigit=removedigit

    def transform(self, X, **transform_params):
        # Nettoyage du texte
        X=X.copy() # pour conserver le fichier d'origine
        return [MyCleanText(text,lowercase=self.lowercase,
            getstemmer=self.getstemmer,
            removestopwords=self.removestopwords,
            getlemmatisation=self.getlemmatisation,
            removedigit=self.removedigit) for text in
X]

```

```

def fit(self, X, y=None, **fit_params):
    return self

```

```

def fit_transform(self, X, y=None, **fit_params):
    return self.fit(X).transform(X)

```

```

def get_params(self, deep=True):
    return {
        'lowercase':self.lowercase,
        'getstemmer':self.getstemmer,
        'removestopwords':self.removestopwords,
        'getlemmatisation':self.getlemmatisation,
        'removedigit':self.removedigit
    }

```

```

def set_params (self, **parameters):

```

```

    for parameter, value in parameters.items():
        setattr(self,parameter,value)
    return self

```

##Etape 1 : Préparer les données

- Load et preparer les données à partir des 2 fichiers csv
- Sélectionner que les lignes où on a True, False

#Ici je cherche à sélectionner que les labels TRUE et FALSE, donc les LIGNES qui contiennent au rating TRUE et FALSE uniquement, le reste on enlève

```

dftrain = pd.read_csv("/content/drive/MyDrive/newsTrain2.csv",
names=['id','text','title','rating'], header=0,sep=',',
encoding='utf8')
dftrain.reset_index(drop = True, inplace = True)

```

```

dftrain2 = pd.read_csv("/content/drive/MyDrive/newsTrain_
_newsTrain.csv", names=['id','text','title','rating'],
header=0,sep=',', encoding='utf8')
dftrain2.reset_index(drop = True, inplace = True)

```

concaténer les deux dataframes en ajoutant les lignes du deuxième à la fin du premier

```

dftrain = pd.concat([dftrain, dftrain2], ignore_index=True)

```

```

dftrain = dftrain.loc[dftrain['rating'].isin(['TRUE','FALSE'])]
#dftest=pd.read_csv("/content/gdrive/MyDrive/newsTest.csv",names=['id'
,'text','title','rating','veracity'], header=0,sep=',',
encoding='utf8')
print("Echantillon de mon dataset \n")
print(dftrain.sample(n=10))
print("\n")
print("Quelques informations importantes \n")
dftrain.info()
#print(dftrain.head())
X_text=dftrain.iloc[0:,1:2]

print("le type de X_test est" ,X_text.columns)
X_title=dftrain.iloc[0:,2:3]
print("le texte est")
display(X_text)
print("le titre est")
display(X_title)
#X_test=dftrain.iloc[1:, :4]
y=dftrain.iloc[0:,-1]
print("voici la dernière case")

```

```

display(y)
#y_test=dftest.iloc[1:, -1]
#y = y_train.ravel()
print("la taille de X_text est",X_text.shape)
print("la taille de y_train est " ,y.shape)
print("les valeurs de TRUE et FALSE sont " ,y.value_counts())

```

Echantillon de mon dataset

	id	text \
212	4celaf1d	Martin Gugino is a 75-year-old professional ag...
2250	d1741354	News "If you won't lead, th...
330	a94e340d	Joseph R. Biden declared last year on the camp...
1289	7c9af097	In Brief The Facts: It is alleged by many in t...
345	0a12bd0d	The NFL has decided to roll back its rule agai...
1043	cef0a3d1	Occupation: President and CEO, Bettencourt Tax...
43	aacdc4d3	Today, Congresswoman Maxine Waters D-CA, Chair...
2315	c9a08752	Tom Clark: Osborne swung his axe without mercy...
71	d0b11b17	163 paid protesters have filed a lawsuit again...
986	d1741354	News "If you won't lead, th...

	title	rating
212	Buffalo Officials Duped By Professional Antifa...	FALSE
2250	IT'S OFFICIAL: Brexit Britain WILL thrive out ...	FALSE
330	Biden's claim about attending historically Bla...	FALSE
1289	Scientist Explains Why He Believes Aluminum Is...	FALSE
345	Pittsburgh Steelers Will Fine ANY Player Kneel...	FALSE
1043	The CDC has admitted face masks do little to p...	TRUE
43	Democratic Lawmaker introduces bill to rename ...	FALSE
2315	Budget 2015: the verdict from our columnists	TRUE
71	163 Paid Protesters Just Filed A Lawsuit Again...	FALSE
986	IT'S OFFICIAL: Brexit Britain WILL thrive out ...	FALSE

Quelques informations importantes

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1578 entries, 0 to 2527
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0   id      1578 non-null    object
1   text    1578 non-null    object
2   title   1554 non-null    object
3   rating  1578 non-null    object
dtypes: object(4)
memory usage: 61.6+ KB
le type de X_test est Index(['text'], dtype='object')
le texte est

```

```

                                text
0    Distracted driving causes more deaths in Canad...
3    But things took a turn for the worse when riot...
4    It's no secret that Epstein and Schiff share a...
6    November 23, 2019 The U.S. Food and Drug Admi...
7    Trump confirms this was a bombing, not an acci...
...
2523 More than four million calls to the taxman are...
2524 More under-18s are being taken to court for se...
2525 The Government's much vaunted Help to Buy Isa ...
2526 The late Robin Williams once called cocaine "G...
2527 The late Robin Williams once called cocaine "G...

```

[1578 rows x 1 columns]

le titre est

```

                                title
0    You Can Be Fined $1,500 If Your Passenger Is U...
3    Obama's Daughters Caught on Camera Burning US ...
4    Leaked Visitor Logs Reveal Schiff's 78 Visits ...
6    FDA Shocking Study: Cells Used In Vaccines Con...
7    Israel Hits Beirut with Nuclear Missile, Trump...
...
2523 Taxman fails to answer four million calls a ye...
2524 Police catch 11-year-olds being used to sell d...
2525 Help to Buy Isa scandal: 500,000 first-time bu...
2526         A coke-snorting generation of hypocrites
2527         A coke-snorting generation of hypocrites

```

[1578 rows x 1 columns]

voici la dernière case

```

0    FALSE
3    FALSE
4    FALSE
6    FALSE
7    FALSE
...
2523    TRUE
2524    TRUE
2525    FALSE
2526    TRUE
2527    TRUE

```

Name: rating, Length: 1578, dtype: object

la taille de X_text est (1578, 1)

la taille de y_train est (1578,)

les valeurs de TRUE et FALSE sont FALSE 1156


```
TRUE          422
Name: rating, dtype: int64
```

Le jeu de données étant déséquilibré, on a pensé à appliquer le downsampling pour équilibrer nos données. On fait en sorte que le nombre de lignes avec en sortie FAUX soit égal au nombre de lignes avec en sortie true.

#On applique du sous-échantillonnage (downsampling) : car on a plus de FALSE (578) que des TRUE (211)

Séparer les classes en deux dataframes

```
df_false = dftrain[dftrain['rating'] == 'FALSE']
df_true = dftrain [dftrain['rating'] == 'TRUE']
```

Sous-échantillonner la classe majoritaire (FALSE) pour obtenir un nombre égal d'échantillons pour chaque classe

```
df_false_subsampled = df_false.sample(n=len(df_true), random_state=42)
```

Concaténer les deux dataframes

```
dftrain = pd.concat([df_false_subsampled, df_true])
```

Mélanger aléatoirement les données

```
dftrain = dftrain.sample(frac=1, random_state=42)
```

```
X_text=dftrain["text"]
X_title=dftrain["title"]
y=dftrain.iloc[:, -1]
print("la taille de X_text est",X_text.shape)
print("\n")
print("la taille de X_title est",X_title.shape)
print("\n")
print("la taille de y_train est " ,y.shape)
print("\n")
print("les valeurs de TRUE et FALSE maintenant sont
" ,y.value_counts())
```

```
la taille de X_text est (844,)
```

```
la taille de X_title est (844,)
```

```
la taille de y_train est  (844,)
```

```
les valeurs de TRUE et FALSE maintenant sont  TRUE          422
```

```
FALSE    422
Name: rating, dtype: int64
```

Découpage du jeu de données en jeu de test et d'entraînement

```
X=dftrain.iloc[0:, 1:3]
print(X)
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size =
0.2,random_state=8)
print("X_train is",X_train.shape)
print("y_train is",y_train.shape)
print("X_test is",X_test.shape)
print("y_test is",y_test.shape)
```

```

                                     text \
615  It's been a long time coming, but finally we h...
1303 Constitutional Attorney Matthew DePerno is an ...
1232 The United States is witnessing a massive, dan...
2022 After three decades on the bench, Sarah Parker...
287  Based on actual results and accounting for sta...
...
1006 5 Million Muslim Children In Yemen Died due to...
1543 The bombshell claim comes from over 20 hours o...
853  BILL GATES EXPLAINS THAT THE COVID VACCINE WIL...
296  Let our journalists help you make sense of the...
1325 Though the whole world relies on RT-PCR to "di...
```

```

                                     title
615  JK Rowling Confirms Stance Against Transgender...
1303 MI Sec of State Official Caught On Video Telli...
1232 What science can tell us about the links betwe...
2022 Sarah Parker leaves legacy on Supreme Court
287  Current Actual Election Result Update: Preside...
...
1006 Re: Meeting the need for isolation space for h...
1543 Breaking: Breonna Taylor's boyfriend says SHE ...
853  A quote from Politifact: Gates never said that...
296  Before This Election, Newt Gingrich Believed t...
1325 COVID19 PCR Tests are Scientifically Meaningless
```

```
[844 rows x 2 columns]
X_train is (675, 2)
y_train is (675,)
X_test is (169, 2)
y_test is (169,)
```

##Etape 2 : Classification selon la colonne TEXT :

Tester avec plusieurs classifieurs classiques.

Ici, c'est une étape importante, on va tester différents classifieurs, pour chacun des classifieurs, on va appliquer le prétraitement + Vectorisation Tfidf, et on applique une cross_val_score avec un Kfold de 10 fois, par la suite on stocke dans une liste all_results la moyenne des accuracy + l'écart type et on la trie par ordre décroissant de moyenne d'accuracy et d'écart type. on remarque que les 3 meilleurs sont SVM,LR et RF qu'on va sélectionner pour leur appliquer le GridSearch sur les paramètres des prétraitements + leurs hyperparamètres pour pouvoir choisir le meilleur.

```
score = 'accuracy'
seed = 7
allresults = []
results = []
names = []
```

```
X_train_text=X_train['text']
X_train_text.reset_index(drop = True, inplace = True)
```

```
# Liste des modèles à tester
```

```
models = [
    ('MultinomialNB', MultinomialNB()),
    ('LogisticRegression', LogisticRegression(random_state=42))
]
```

```
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('RF', RandomForestClassifier()))
models.append(('SVM', SVC()))
```

```
# Création d'un pipeline pour chaque modèle
```

```
pipelines = []
for name,model in models:
    pipeline = Pipeline([
        ('normalize', TextNormalizer()),
        ('tfidf', TfidfVectorizer()),
        (name,model)
    ])
    pipelines.append((name,pipeline))
#pipeline.fit(X_train_text,y_train)
all_results=[]
scores=[]
for p in pipelines:
    print(p[1])
    # cross validation en 10 fois
    kfold = KFold(n_splits=10,random_state=seed,shuffle=True)
```

```

# print ("Evaluation de ",p)
start_time = time.time()
# application de la classification
cv_results = cross_val_score(p[1],X_train_text,y_train, cv=kfold,
scoring=score)
#print("Pour le classifieur",p[0],"on a un score
de",cv_results.mean(),"et un écart type de",cv_results.std())
scores.append(cv_results)
names.append(p[0])

all_results.append((p[0],cv_results.mean(),cv_results.std()))
end_time = time.time()

all_results = sorted(all_results, key=lambda x: (-x[1], -x[2]))
print("all resultats", all_results)
# affichage des résultats
#print ('\nLe meilleur resultat : ',max(results))

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('MultinomialNB', MultinomialNB())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('LogisticRegression',
LogisticRegression(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('KNN', KNeighborsClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('CART', DecisionTreeClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('RF', RandomForestClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('SVM', SVC())])
all resultats [('SVM', 0.9038410886742756, 0.043294516291260216),
('RF', 0.8948200175592624, 0.0370882010820444), ('LogisticRegression',
0.8356233538191397, 0.0345425007546011), ('CART', 0.832550482879719,
0.05427625364367281), ('MultinomialNB', 0.8265144863915715,
0.06346830897936112), ('KNN', 0.7051141352063214,
0.06333815247759318)]

```

On affiche les accuracy de chaque classifieur, on remarque la médiane (en rouge) de chaque et l'écart type aussi.

```

import matplotlib.pyplot as plt
fig = plt.figure()


```

```

fig.suptitle('Comparaison des algorithmes')
ax = fig.add_subplot(111)
plt.boxplot(scores)
ax.set_xticklabels(names)

[Text(1, 0, 'MultinomialNB'),
 Text(2, 0, 'LogisticRegression'),
 Text(3, 0, 'KNN'),
 Text(4, 0, 'CART'),
 Text(5, 0, 'RF'),
 Text(6, 0, 'SVM')]

```



Choisir les meilleurs paramètres et hyperparamètres pour SVM et RF :

On a un pipeline pour chaque prétraitement différent, on essaye pas mal (miniscule, lemmatisation, miniscule + lemmatisation..) et on stocke le fit_transform de nos X_train, X_test sur les pipelines dans des listes qui vont contenir tous les fit_transform des pipelines pour chaque classifieur, par la suite on parcourt ces listes là, on itère dessus, et chaque élément de la liste (train) va passer par le GridSearch et puis on prédit sur son correspondant dans liste (test).

```

from sklearn.model_selection import GridSearchCV

```

```

from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from tabulate import tabulate
import numpy as np
np.random.seed(42) # Set the random seed for NumPy
print("y_train", y_train.shape)
print("y_test", y_test.shape)
print("X_test", X_test.shape)

X_train_text=X_train['text']
X_train_title=X_train['title']

# le plus simple est de faire un test sur differents pipelines.
# pipeline de l'utilisation de CountVectorizer sur le texte avec
differents pre-traitements
CV_brut = Pipeline([('cleaner', TextNormalizer()),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])
CV_lowcase = Pipeline([('cleaner',
                        TextNormalizer(removestopwords=False, lowercase=True,

getstemmer=False, removedigit=False)),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])
CV_lowStop = Pipeline([('cleaner',
                        TextNormalizer(removestopwords=True, lowercase=True,

getstemmer=False, removedigit=False)),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])

CV_lowStopstem = Pipeline([('cleaner',
                            TextNormalizer(removestopwords=True, lowercase=True,

getstemmer=True, removedigit=False)),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])

# pipeline de l'utilisation de TfidfVectorizer avec differents pre-
traitements

```

```
TFIDF_brut = Pipeline ([('cleaner', TextNormalizer()),
                        ('tfidf_vectorizer',
                         TfidfVectorizer(lowercase=False))])
```

```
TFIDF_lowercase = Pipeline([('cleaner',
                             TextNormalizer(removestopwords=False, lowercase=True,
```

```
                             getstemmer=False, removedigit=False)),
                          ('tfidf_vectorizer',
                           TfidfVectorizer(lowercase=False))])
```

```
TFIDF_lowStop = Pipeline([('cleaner',
                           TextNormalizer(removestopwords=True, lowercase=True,
```

```
                           getstemmer=False, removedigit=False)),
                          ('tfidf_vectorizer',
                           TfidfVectorizer(lowercase=False))])
```

```
TFIDF_lowStopstem = Pipeline([('cleaner',
                               TextNormalizer(removestopwords=True, lowercase=True,
```

```
                               getstemmer=True, removedigit=False)),
                              ('tfidf_vectorizer',
                               TfidfVectorizer(lowercase=False))])
```

```
# Liste de tous les modeles à tester
```

```
all_models = [
    ("CV_brut", CV_brut),
    ("CV_lowercase", CV_lowercase),
    ("CV_lowStop", CV_lowStop),
    ("CV_lowStopstem", CV_lowStopstem),
    ("TFIDF_lowercase", TFIDF_lowercase),
    ("TFIDF_lowStop", TFIDF_lowStop),
    ("TFIDF_lowStopstem", TFIDF_lowStopstem),
    ("TFIDF_brut", TFIDF_brut)
]
```

```
X_train_text_SVC = []
X_test_text_SVC = []
```

```
X_train_text_RandomForestClassifier = []
X_test_text_RandomForestClassifier = []
```

```
for name, pipeline in all_models :
```

```
X_train_text_SVC.append(pipeline.fit_transform(X_train['text']).toarray())
```

```

X_test_text_SVC.append(pipeline.transform(X_test['text']).toarray())

X_train_text_RandomForestClassifier.append(pipeline.fit_transform(X_train['text']).toarray())

X_test_text_RandomForestClassifier.append(pipeline.transform(X_test['text']).toarray())


models = {
    'SVC': SVC(random_state=42),
    'RandomForestClassifier': RandomForestClassifier(random_state=42)
}

params = {'SVC': [{'C': [0.001, 0.01, 0.1, 1, 2, 5, 7, 10]},
                  {'gamma': [0.001, 0.01, 0.1, 0.2, 0.3, 0.5, 0.7, 1]}],
          'RandomForestClassifier': [{'n_estimators': [10, 50, 100, 200, 300]},
                                     {'max_features': ['auto', 'sqrt', 'log2']}]}

for model_name, model in models.items():
    score='accuracy'
    X_train_text = eval('X_train_text_' + model_name)
    X_test_text = eval('X_test_text_' + model_name)
    for i in range(len(X_train_text)):
        grid_search = GridSearchCV(model, params[model_name], n_jobs=-1,
        verbose=1, scoring=score)
        print("grid search fait")
        print("X_train",X_train_text[i].shape)
        print("y_train",y_train.shape)
        grid_search.fit(X_train_text[i],y_train)
        print ('meilleur score %0.3f'%(grid_search.best_score_),'\n')
        print ('meilleur estimateur',grid_search.best_estimator_,'\n')
        y_pred = grid_search.predict(X_test_text[i])
        MyshowAllScores(y_test,y_pred)

        print("Ensemble des meilleurs paramètres :")
        best_parameters = grid_search.best_estimator_.get_params()
        for param_dict in params[model_name]:
            for param_name, param_value in param_dict.items():
                print("\t%s: %r" % (param_name,
                best_parameters[param_name]))

```



```
y_train (675,)
y_test (169,)
X_test (169, 2)
grid search fait
X_train (675, 24296)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.849
```

meilleur estimateur SVC(kernel='linear', random_state=42)

Accuracy : 0.870

Classification Report

	precision	recall	f1-score	support
FALSE	0.89610	0.83133	0.86250	83
TRUE	0.84783	0.90698	0.87640	86
accuracy			0.86982	169
macro avg	0.87196	0.86915	0.86945	169
weighted avg	0.87154	0.86982	0.86958	169

Ensemble des meilleurs paramètres :

```
C: 1.0
gamma: 'scale'
kernel: 'linear'
grid search fait
X_train (675, 20741)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.846
```

meilleur estimateur SVC(gamma=0.01, random_state=42)

Accuracy : 0.905

Classification Report

	precision	recall	f1-score	support
FALSE	0.83838	1.00000	0.91209	83
TRUE	1.00000	0.81395	0.89744	86
accuracy			0.90533	169
macro avg	0.91919	0.90698	0.90476	169
weighted avg	0.92063	0.90533	0.90463	169

Ensemble des meilleurs paramètres :

```
C: 1.0
gamma: 0.01
kernel: 'rbf'
```

```
grid search fait
X_train (675, 20601)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.852
```

meilleur estimateur SVC(C=10, random_state=42)

Accuracy : 0.882

Classification Report

	precision	recall	f1-score	support
FALSE	0.90909	0.84337	0.87500	83
TRUE	0.85870	0.91860	0.88764	86
accuracy			0.88166	169
macro avg	0.88389	0.88099	0.88132	169
weighted avg	0.88345	0.88166	0.88143	169

Ensemble des meilleurs paramètres :

C: 10

gamma: 'scale'

kernel: 'rbf'

```
grid search fait
```

```
X_train (675, 14012)
```

```
y_train (675,)
```

```
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.856
```

meilleur estimateur SVC(C=10, random_state=42)

Accuracy : 0.882

Classification Report

	precision	recall	f1-score	support
FALSE	0.88889	0.86747	0.87805	83
TRUE	0.87500	0.89535	0.88506	86
accuracy			0.88166	169
macro avg	0.88194	0.88141	0.88155	169
weighted avg	0.88182	0.88166	0.88162	169

Ensemble des meilleurs paramètres :

C: 10

gamma: 'scale'

kernel: 'rbf'

```
grid search fait
```

```
X_train (675, 20741)
```

```
y_train (675,)
```

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.886

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.911

Classification Report

	precision	recall	f1-score	support
FALSE	0.92500	0.89157	0.90798	83
TRUE	0.89888	0.93023	0.91429	86
accuracy			0.91124	169
macro avg	0.91194	0.91090	0.91113	169
weighted avg	0.91171	0.91124	0.91119	169

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

X_train (675, 20601)

y_train (675,)

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.887

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.917

Classification Report

	precision	recall	f1-score	support
FALSE	0.88764	0.95181	0.91860	83
TRUE	0.95000	0.88372	0.91566	86
accuracy			0.91716	169
macro avg	0.91882	0.91776	0.91713	169
weighted avg	0.91937	0.91716	0.91711	169

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

X_train (675, 14012)

y_train (675,)

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.887

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.929

Classification Report

	precision	recall	f1-score	support
FALSE	0.90805	0.95181	0.92941	83
TRUE	0.95122	0.90698	0.92857	86
accuracy			0.92899	169
macro avg	0.92963	0.92939	0.92899	169
weighted avg	0.93002	0.92899	0.92898	169

Ensemble des meilleurs paramètres :

C: 1
gamma: 'scale'
kernel: 'rbf'
grid search fait
X_train (675, 24296)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.886

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.899

Classification Report

	precision	recall	f1-score	support
FALSE	0.90244	0.89157	0.89697	83
TRUE	0.89655	0.90698	0.90173	86
accuracy			0.89941	169
macro avg	0.89950	0.89927	0.89935	169
weighted avg	0.89944	0.89941	0.89939	169

Ensemble des meilleurs paramètres :

C: 2
gamma: 'scale'
kernel: 'rbf'
grid search fait
X_train (675, 24296)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.873

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.882

Classification Report

	precision	recall	f1-score	support
FALSE	0.87952	0.87952	0.87952	83
TRUE	0.88372	0.88372	0.88372	86
accuracy			0.88166	169
macro avg	0.88162	0.88162	0.88162	169
weighted avg	0.88166	0.88166	0.88166	169

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'sqrt'

grid search fait

X_train (675, 20741)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.867

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.864

Classification Report

	precision	recall	f1-score	support
FALSE	0.87500	0.84337	0.85890	83
TRUE	0.85393	0.88372	0.86857	86
accuracy			0.86391	169
macro avg	0.86447	0.86355	0.86373	169
weighted avg	0.86428	0.86391	0.86382	169

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

X_train (675, 20601)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.868

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.893

Classification Report

	precision	recall	f1-score	support
FALSE	0.90123	0.87952	0.89024	83

TRUE	0.88636	0.90698	0.89655	86
accuracy			0.89349	169
macro avg	0.89380	0.89325	0.89340	169
weighted avg	0.89367	0.89349	0.89345	169

Ensemble des meilleurs paramètres :

```

n_estimators: 100
max_features: 'sqrt'
grid search fait
X_train (675, 14012)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.867

```

meilleur estimateur RandomForestClassifier(max_features='log2', random_state=42)

Accuracy : 0.917

Classification Report

	precision	recall	f1-score	support
FALSE	0.90588	0.92771	0.91667	83
TRUE	0.92857	0.90698	0.91765	86
accuracy			0.91716	169
macro avg	0.91723	0.91734	0.91716	169
weighted avg	0.91743	0.91716	0.91717	169

Ensemble des meilleurs paramètres :

```

n_estimators: 100
max_features: 'log2'
grid search fait
X_train (675, 20741)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.879

```

meilleur estimateur RandomForestClassifier(n_estimators=300, random_state=42)

Accuracy : 0.911

Classification Report

	precision	recall	f1-score	support
FALSE	0.90476	0.91566	0.91018	83
TRUE	0.91765	0.90698	0.91228	86
accuracy			0.91124	169

macro avg	0.91120	0.91132	0.91123	169
weighted avg	0.91132	0.91124	0.91125	169

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

X_train (675, 20601)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.879

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.917

Classification Report

	precision	recall	f1-score	support
FALSE	0.92593	0.90361	0.91463	83
TRUE	0.90909	0.93023	0.91954	86
accuracy			0.91716	169
macro avg	0.91751	0.91692	0.91709	169
weighted avg	0.91736	0.91716	0.91713	169

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'sqrt'

grid search fait

X_train (675, 14012)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.877

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.911

Classification Report

	precision	recall	f1-score	support
FALSE	0.90476	0.91566	0.91018	83
TRUE	0.91765	0.90698	0.91228	86
accuracy			0.91124	169
macro avg	0.91120	0.91132	0.91123	169
weighted avg	0.91132	0.91124	0.91125	169

Ensemble des meilleurs paramètres :

```
n_estimators: 300
max_features: 'sqrt'
grid search fait
X_train (675, 24296)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.880
```

meilleur estimateur RandomForestClassifier(random_state=42)


Accuracy : 0.905



Classification Report



	precision	recall	f1-score	support
FALSE	0.90361	0.90361	0.90361	83
TRUE	0.90698	0.90698	0.90698	86
accuracy			0.90533	169
macro avg	0.90530	0.90530	0.90530	169
weighted avg	0.90533	0.90533	0.90533	169


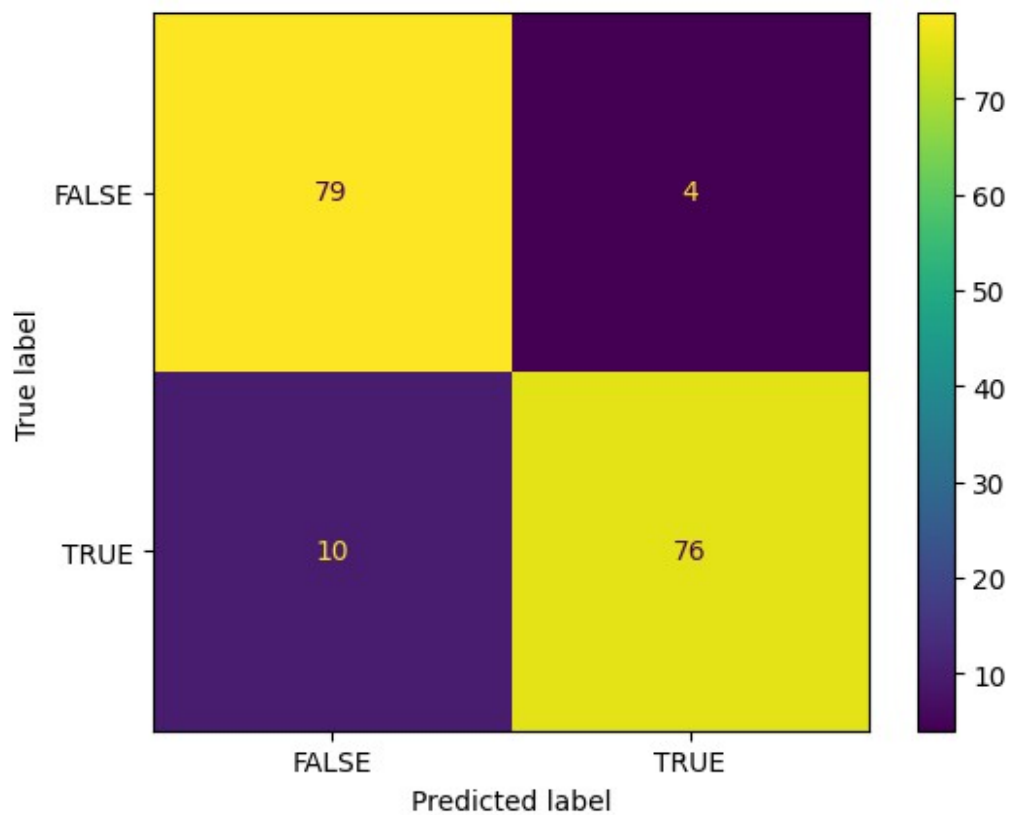
Ensemble des meilleurs paramètres :



```
n_estimators: 100
max_features: 'sqrt'
```


























##Etape 3 : Classification selon la colonne TITRE :

```

from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
import time
import numpy as np

# Utilisez la méthode ravel() pour transformer y_train en un tableau
# unidimensionnel
#X_train = np.ravel(X_train)
print("X_train", X_train.shape)
print("y_train", y_train.shape)
np.random.seed(42) # Set the random seed for NumPy

score = 'accuracy'
seed = 7
allresults = []

```



```

results = []
names = []

# Liste des modèles à tester
models = [
    ('MultinomialNB', MultinomialNB()),
    ('LogisticRegression', LogisticRegression(random_state=42))
]

models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier(random_state=42)))
models.append(('RF', RandomForestClassifier(random_state=42)))
models.append(('SVM', SVC(random_state=42)))

# Création d'un pipeline pour chaque modèle
pipelines = []
for name,model in models:
    pipeline = Pipeline([
        ('normalize', TextNormalizer()),
        ('tfidf', TfidfVectorizer()),
        (name,model)
    ])
    pipelines.append((name,pipeline))
    #pipeline.fit(X_train,y_train)
all_results=[]
scores=[]
names = []
for p in pipelines:
    print(p[1])
    # cross validation en 10 fois
    kfold = KFold(n_splits=10,random_state=seed,shuffle=True)

    # print ("Evaluation de ",p)
    start_time = time.time()
    # application de la classification
    cv_results = cross_val_score(p[1],X_train['title'],y_train,
cv=kfold, scoring=score)
    #print("Pour le classifieur",p[0],"on a un score
de",cv_results.mean(),"et un écart type de",cv_results.std())
    scores.append(cv_results)
    names.append(p[0])
    all_results.append((p[0],cv_results.mean(),cv_results.std()))
    end_time = time.time()

print("all resultats", all_results)

```

```

all_results = sorted(all_results, key=lambda x: (-x[1], -x[2]))
print("all resultats", all_results)
    # affichage des résultats
#print ('\nLe meilleur resultat : ',max(results))

X_train (675, 2)
y_train (675,)
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('MultinomialNB', MultinomialNB())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('LogisticRegression',
LogisticRegression(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('KNN', KNeighborsClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('CART', DecisionTreeClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('RF', RandomForestClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('SVM', SVC(random_state=42))])
all resultats [('MultinomialNB', 0.810359964881475,
0.05872410123263959), ('LogisticRegression', 0.8104697102721685,
0.07018958627295246), ('KNN', 0.6446224758560141,
0.07936226071732319), ('CART', 0.8104697102721685,
0.07035617709825927), ('RF', 0.8416154521510096, 0.0799660194065715),
('SVM', 0.8667251975417033, 0.08008579483575896)]
all resultats [('SVM', 0.8667251975417033, 0.08008579483575896),
('RF', 0.8416154521510096, 0.0799660194065715), ('CART',
0.8104697102721685, 0.07035617709825927), ('LogisticRegression',
0.8104697102721685, 0.07018958627295246), ('MultinomialNB',
0.810359964881475, 0.05872410123263959), ('KNN', 0.6446224758560141,
0.07936226071732319)]

```

On affiche les accuracy de chaque classifieur, on remarque la médiane (en rouge) de chaque et l'écart type aussi.

```


import matplotlib.pyplot as plt
fig = plt.figure()
fig.suptitle('Comparaison des algorithmes')
ax = fig.add_subplot(111)
plt.boxplot(scores)
ax.set_xticklabels(names)

[Text(1, 0, 'MultinomialNB'),
Text(2, 0, 'LogisticRegression'),

```

```
Text(3, 0, 'KNN'),
Text(4, 0, 'CART'),
Text(5, 0, 'RF'),
Text(6, 0, 'SVM')]
```

Comparaison des algorithmes



Choisir les meilleurs paramètres pour SVM et RF :

On a un pipeline pour chaque prétraitement différent, on essaye pas mal (miniscule, lemmatisation, miniscule + lemmatisation..) et on stocke le fit_transform de nos X_train, X_test sur les pipelines dans des listes qui vont contenir tous les fit_transform des pipelines pour chaque classifieur, par la suite on parcourt ces listes là, on itère dessus, et chaque élément de la liste (train) va passer par le GridSearch et puis on prédit sur son correspondant dans la liste (test).

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
```

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from tabulate import tabulate
import numpy as np
np.random.seed(42) # Set the random seed for NumPy
print("y_train", y_train.shape)
print("y_test", y_test.shape)
#X_test = np.ravel(X_test)
print("X_test", X_test.shape)
# le plus simple est de faire un test sur differents pipelines.
# pipeline de l'utilisation de CountVectorizer sur le texte avec
differents pre-traitements
CV_brut = Pipeline([('cleaner', TextNormalizer()),
                    ('count_vectorizer',
CountVectorizer(lowercase=False))])
CV_lowercase = Pipeline([('cleaner',
TextNormalizer(removestopwords=False, lowercase=True,
getstemmer=False, removedigit=False)),
                        ('count_vectorizer',
CountVectorizer(lowercase=False))])
CV_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True, lowercase=True,
getstemmer=False, removedigit=False)),
                      ('count_vectorizer',
CountVectorizer(lowercase=False))])
CV_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True, lowercase=True,
getstemmer=True, removedigit=False)),
                          ('count_vectorizer',
CountVectorizer(lowercase=False))])
# pipeline de l'utilisation de TfidfVectorizer avec differents pre-
traitements
TFIDF_brut = Pipeline ([('cleaner', TextNormalizer()),
                       ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])
TFIDF_lowercase = Pipeline([('cleaner',
TextNormalizer(removestopwords=False, lowercase=True,
getstemmer=False, removedigit=False)),
                           ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])
```

```

TFIDF_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
                        ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

```

```

TFIDF_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
                        ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

```

```

# Liste de tous les modeles à tester
all_models = [
    ("CV_brut", CV_brut),
    ("CV_lowcase", CV_lowcase),
    ("CV_lowStop", CV_lowStop),
    ("CV_lowStopstem", CV_lowStopstem),
    ("TFIDF_lowcase", TFIDF_lowcase),
    ("TFIDF_lowStop", TFIDF_lowStop),
    ("TFIDF_lowStopstem", TFIDF_lowStopstem),
    ("TFIDF_brut", TFIDF_brut)
]

```

```

X_train_title_SVC = []
X_test_title_SVC = []

```

```

X_train_title_RandomForestClassifier = []
X_test_title_RandomForestClassifier = []

```

```

for name, pipeline in all_models :

```

```

    X_train_title_SVC.append(pipeline.fit_transform(X_train['title']).toar
ray())

```

```

    X_test_title_SVC.append(pipeline.transform(X_test['title']).toarray())

```

```

    X_train_title_RandomForestClassifier.append(pipeline.fit_transform(X_t
rain['title']).toarray())

```

```

    X_test_title_RandomForestClassifier.append(pipeline.transform(X_test['
title']).toarray())

```

```

models = {
    'SVC': SVC(random_state=42),
    'RandomForestClassifier': RandomForestClassifier(random_state=42)
}

params = {'SVC': [{'C': [0.001, 0.01, 0.1, 1, 2, 5, 7, 10]},
                  {'gamma': [0.001, 0.01, 0.1, 0.2, 0.3, 0.5, 0.7, 1]}],
          {'kernel': ['linear', 'rbf']}],
          'RandomForestClassifier': [{'n_estimators': [10, 50, 100, 200,
300]}],
                                     {'max_features': ['auto', 'sqrt',
'log2']}]
}

for model_name, model in models.items():
    score='accuracy'
    X_train_title = eval('X_train_title_' + model_name)
    X_test_title = eval('X_test_title_' + model_name)
    for i in range(len(X_train_title)):
        grid_search = GridSearchCV(model, params[model_name], n_jobs=-1,
verbose=1, scoring=score)
        print("grid search fait")
        print("X_train", X_train_title[i].shape)
        print("y_train", y_train.shape)
        grid_search.fit(X_train_title[i], y_train)
        print('meilleur score %0.3f'%(grid_search.best_score_), '\n')
        print('meilleur estimateur', grid_search.best_estimator_, '\n')
        y_pred = grid_search.predict(X_test_title[i])
        MyshowAllScores(y_test, y_pred)

        print("Ensemble des meilleurs paramètres :")
        best_parameters = grid_search.best_estimator_.get_params()
        for param_dict in params[model_name]:
            for param_name, param_value in param_dict.items():
                print("\t%s: %r" % (param_name,
best_parameters[param_name]))

y_train (675,)
y_test (169,)
X_test (169, 2)
grid search fait
X_train (675, 5734)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.831

meilleur estimateur SVC(gamma=0.3, random_state=42)

```

Accuracy : 0.882

Classification Report

	precision	recall	f1-score	support
FALSE	0.89873	0.85542	0.87654	83
TRUE	0.86667	0.90698	0.88636	86
accuracy			0.88166	169
macro avg	0.88270	0.88120	0.88145	169
weighted avg	0.88242	0.88166	0.88154	169

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.3

kernel: 'rbf'

grid search fait

X_train (675, 4861)

y_train (675,)

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.840

meilleur estimateur SVC(gamma=0.3, random_state=42)

Accuracy : 0.917

Classification Report

	precision	recall	f1-score	support
FALSE	0.92593	0.90361	0.91463	83
TRUE	0.90909	0.93023	0.91954	86
accuracy			0.91716	169
macro avg	0.91751	0.91692	0.91709	169
weighted avg	0.91736	0.91716	0.91713	169

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.3

kernel: 'rbf'

grid search fait

X_train (675, 4733)

y_train (675,)

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.831

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.911

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

FALSE	0.94737	0.86747	0.90566	83
TRUE	0.88172	0.95349	0.91620	86
accuracy			0.91124	169
macro avg	0.91454	0.91048	0.91093	169
weighted avg	0.91396	0.91124	0.91102	169

Ensemble des meilleurs paramètres :

```

C: 1.0
gamma: 0.5
kernel: 'rbf'
grid search fait
X_train (675, 3613)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.833

```

meilleur estimateur SVC(gamma=0.7, random_state=42)

Accuracy : 0.905

Classification Report

	precision	recall	f1-score	support
FALSE	0.86813	0.95181	0.90805	83
TRUE	0.94872	0.86047	0.90244	86
accuracy			0.90533	169
macro avg	0.90842	0.90614	0.90524	169
weighted avg	0.90914	0.90533	0.90519	169

Ensemble des meilleurs paramètres :

```

C: 1.0
gamma: 0.7
kernel: 'rbf'
grid search fait
X_train (675, 4861)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.834

```

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.888

Classification Report

	precision	recall	f1-score	support
FALSE	0.88095	0.89157	0.88623	83
TRUE	0.89412	0.88372	0.88889	86

accuracy			0.88757	169
macro avg	0.88754	0.88764	0.88756	169
weighted avg	0.88765	0.88757	0.88758	169

Ensemble des meilleurs paramètres :

```

C: 1
gamma: 'scale'
kernel: 'rbf'
grid search fait
X_train (675, 4733)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.836

```

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.888

Classification Report

	precision	recall	f1-score	support
FALSE	0.88095	0.89157	0.88623	83
TRUE	0.89412	0.88372	0.88889	86

accuracy			0.88757	169
macro avg	0.88754	0.88764	0.88756	169
weighted avg	0.88765	0.88757	0.88758	169

Ensemble des meilleurs paramètres :

```

C: 1
gamma: 'scale'
kernel: 'rbf'
grid search fait
X_train (675, 3613)
y_train (675,)
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.830

```

meilleur estimateur SVC(gamma=1, random_state=42)

Accuracy : 0.893

Classification Report

	precision	recall	f1-score	support
FALSE	0.88235	0.90361	0.89286	83
TRUE	0.90476	0.88372	0.89412	86

accuracy			0.89349	169
macro avg	0.89356	0.89367	0.89349	169

weighted avg 0.89376 0.89349 0.89350 169

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 1

kernel: 'rbf'

grid search fait

X_train (675, 5734)

y_train (675,)

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.837

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.893

Classification Report

	precision	recall	f1-score	support
FALSE	0.88235	0.90361	0.89286	83
TRUE	0.90476	0.88372	0.89412	86
accuracy			0.89349	169
macro avg	0.89356	0.89367	0.89349	169
weighted avg	0.89376	0.89349	0.89350	169

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

X_train (675, 5734)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.824

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.911

Classification Report

	precision	recall	f1-score	support
FALSE	0.90476	0.91566	0.91018	83
TRUE	0.91765	0.90698	0.91228	86
accuracy			0.91124	169
macro avg	0.91120	0.91132	0.91123	169
weighted avg	0.91132	0.91124	0.91125	169

Ensemble des meilleurs paramètres :
n_estimators: 50
max_features: 'sqrt'
grid search fait
X_train (675, 4861)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.824

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.888

Classification Report

	precision	recall	f1-score	support
FALSE	0.90000	0.86747	0.88344	83
TRUE	0.87640	0.90698	0.89143	86
accuracy			0.88757	169
macro avg	0.88820	0.88722	0.88743	169
weighted avg	0.88799	0.88757	0.88750	169

Ensemble des meilleurs paramètres :
n_estimators: 100
max_features: 'sqrt'
grid search fait
X_train (675, 4733)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.827

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.899

Classification Report

	precision	recall	f1-score	support
FALSE	0.88372	0.91566	0.89941	83
TRUE	0.91566	0.88372	0.89941	86
accuracy			0.89941	169
macro avg	0.89969	0.89969	0.89941	169
weighted avg	0.89998	0.89941	0.89941	169

Ensemble des meilleurs paramètres :
n_estimators: 50
max_features: 'sqrt'
grid search fait

```
X_train (675, 3613)
y_train (675,)
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.822
```

```
meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)
```

Accuracy : 0.876

Classification Report

	precision	recall	f1-score	support
FALSE	0.86047	0.89157	0.87574	83
TRUE	0.89157	0.86047	0.87574	86
accuracy			0.87574	169
macro avg	0.87602	0.87602	0.87574	169
weighted avg	0.87629	0.87574	0.87574	169

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

```
X_train (675, 4861)
```

```
y_train (675,)
```

```
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.831
```

```
meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)
```

Accuracy : 0.858

Classification Report

	precision	recall	f1-score	support
FALSE	0.85542	0.85542	0.85542	83
TRUE	0.86047	0.86047	0.86047	86
accuracy			0.85799	169
macro avg	0.85794	0.85794	0.85794	169
weighted avg	0.85799	0.85799	0.85799	169

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

```
X_train (675, 4733)
```

```
y_train (675,)
```

```
Fitting 5 folds for each of 8 candidates, totalling 40 fits
```

meilleur score 0.825

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.876

Classification Report

	precision	recall	f1-score	support
FALSE	0.86047	0.89157	0.87574	83
TRUE	0.89157	0.86047	0.87574	86
accuracy			0.87574	169
macro avg	0.87602	0.87602	0.87574	169
weighted avg	0.87629	0.87574	0.87574	169

Ensemble des meilleurs paramètres :

n_estimators: 50

max_features: 'sqrt'

grid search fait

X_train (675, 3613)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.824

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.870

Classification Report

	precision	recall	f1-score	support
FALSE	0.85882	0.87952	0.86905	83
TRUE	0.88095	0.86047	0.87059	86
accuracy			0.86982	169
macro avg	0.86989	0.86999	0.86982	169
weighted avg	0.87008	0.86982	0.86983	169

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'sqrt'

grid search fait

X_train (675, 5734)

y_train (675,)

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.830

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.870


Classification Report



	precision	recall	f1-score	support
FALSE	0.85882	0.87952	0.86905	83
TRUE	0.88095	0.86047	0.87059	86
accuracy			0.86982	169
macro avg	0.86989	0.86999	0.86982	169
weighted avg	0.87008	0.86982	0.86983	169



Ensemble des meilleurs paramètres :



n_estimators: 50



max_features: 'sqrt'


























##Etape 4 : Classification selon le TITRE ET TEXT ENSEMBLE (Concaténés):

- On va à partir de X_train concaténer les 2 colonnes TEXT et TITLE en mettant un espace entre les deux
- Vu qu'on va travailler sur la colonne text_titre qu'on vient de créer, on va sélectionner cette dernière depuis le X_train et X_test pour apprendre et tester après.

#concaténation

```
X_train=X_train.apply(lambda row: ' '.join([str(val) for val in row]),
axis=1)
X_test=X_test.apply(lambda row: ' '.join([str(val) for val in row]),
axis=1)
```

Ici, c'est une étape importante, on va tester différents classifieurs, pour chacun des classifieurs, on va appliquer le prétraitement + Vectorisation Tfidf, et on applique une cross_val_score avec un Kfold de 10 fois, par la suite on stocke dans une liste all_results la moyenne des accuracy + l'écart type et on la trie par ordre décroissant de moyenne d'accuracy et d'écart type. on remarque que les 2 meilleurs sont SVM et RF qu'on va sélectionner pour leur appliquer le GridSearch sur les paramètres des prétraitements + leurs hyperparamètres pour pouvoir choisir le meilleur.

```
from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
import time
import numpy as np

# Utilisez la méthode ravel() pour transformer y_train en un tableau
unidimensionnel
y_train = np.ravel(y_train)

np.random.seed(42) # Set the random seed for NumPy

score = 'accuracy'
seed = 7
allresults = []
results = []
names = []

# Liste des modèles à tester
models = [
    ('MultinomialNB', MultinomialNB()),
    ('LogisticRegression', LogisticRegression(random_state=42))
]

#models.append(('LR', LogisticRegression(solver='lbfgs')))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier(random_state=42)))
models.append(('RF', RandomForestClassifier(random_state=42)))
models.append(('SVM', SVC(random_state=42)))

# Création d'un pipeline pour chaque modèle
pipelines = []
for name,model in models:
    pipeline = Pipeline([
        ('normalize', TextNormalizer()),
        ('tfidf', TfidfVectorizer()),
        (name,model)
    ])
    pipelines.append((name,pipeline))

all_results=[]
scores=[]

```

```

for p in pipelines:
    print(p[1])
    # cross validation en 10 fois
    kfold = KFold(n_splits=10, random_state=seed, shuffle=True)

    # print ("Evaluation de ",p)
    start_time = time.time()
    # application de la classification
    cv_results = cross_val_score(p[1],X_train,y_train, cv=kfold,
    scoring=score)
    #print("Pour le classifieur",p[0],"on a un score
    de",cv_results.mean(),"et un écart type de",cv_results.std())
    scores.append(cv_results)

    all_results.append((p[0],cv_results.mean(),cv_results.std()))
    end_time = time.time()

print("all resultats", all_results)

all_results = sorted(all_results, key=lambda x: (-x[1], -x[2]))
print("all resultats", all_results)
    # affichage des résultats
    #print ('\nLe meilleur resultat : ',max(results))

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('MultinomialNB', MultinomialNB())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('LogisticRegression',
LogisticRegression(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('KNN', KNeighborsClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('CART', DecisionTreeClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('RF', RandomForestClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('SVM', SVC(random_state=42))])
all resultats [('MultinomialNB', 0.8041703248463564,
0.07914308005639414), ('LogisticRegression', 0.8267778753292362,
0.028218150067996238), ('KNN', 0.6814310798946445,
0.07114977093131233), ('CART', 0.8118305531167691,
0.039914873843984094), ('RF', 0.891812993854258,
0.038760785601535896), ('SVM', 0.8949517120280948,

```



```

CV_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
                          ('count_vectorizer',
CountVectorizer(lowercase=False))])

# pipeline de l'utilisation de TfidfVectorizer avec differents pre-
traitements
TFIDF_brut = Pipeline ([('cleaner', TextNormalizer()),
                        ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowercase = Pipeline([('cleaner',
TextNormalizer(removestopwords=False,lowercase=True,

getstemmer=False,removedigit=False)),
                          ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])
TFIDF_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
                          ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
                          ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

# Liste de tous les modeles à tester
all_models = [
    ("CV_brut", CV_brut),
    ("CV_lowercase", CV_lowercase),
    ("CV_lowStop", CV_lowStop),
    ("CV_lowStopstem", CV_lowStopstem),
    ("TFIDF_lowercase", TFIDF_lowercase),
    ("TFIDF_lowStop", TFIDF_lowStop),
    ("TFIDF_lowStopstem", TFIDF_lowStopstem),
    ("TFIDF_brut", TFIDF_brut)
]

X_train_text_title_SVC = []
X_test_text_title_SVC = []

```

```

X_train_text_title_RandomForestClassifier = []
X_test_text_title_RandomForestClassifier = []

for name, pipeline in all_models :

X_train_text_title_SVC.append(pipeline.fit_transform(X_train).toarray(
))
    X_test_text_title_SVC.append(pipeline.transform(X_test).toarray())

X_train_text_title_RandomForestClassifier.append(pipeline.fit_transfor
m(X_train).toarray())

X_test_text_title_RandomForestClassifier.append(pipeline.transform(X_t
est).toarray())

```

```

models = {
    'SVC': SVC(random_state=42),
    'RandomForestClassifier': RandomForestClassifier(random_state=42)
}

```

```

params = {'SVC': [{'C': [0.001, 0.01, 0.1, 1,2,5,7,10]},
                  {'gamma': [0.001, 0.01, 0.1,0.2,0.3,0.5,0.7,1]}],
          {'kernel': ['linear', 'rbf']}],
    'RandomForestClassifier': [{'n_estimators': [10, 50, 100, 200,
300]}],
                                {'max_features': ['auto', 'sqrt',
'log2']}]],
}

```

```

for model_name, model in models.items():
    score='accuracy'
    X_train_text_title = eval('X_train_text_' + model_name)
    X_test_text_title = eval('X_test_text_' + model_name)
    for i in range (len(X_train_text_title)):
        grid_search = GridSearchCV(model, params[model_name], n_jobs=-1,
verbose=1,scoring=score)
        print("grid search fait")
        grid_search.fit(X_train_text_title[i],y_train)
        print ('meilleur score %.3f'%(grid_search.best_score_),'\n')
        print ('meilleur estimateur',grid_search.best_estimator_,'\n')
        y_pred = grid_search.predict(X_test_text_title[i])
        MyshowAllScores(y_test,y_pred)

    print("Ensemble des meilleurs paramètres :")

```

```

        best_parameters = grid_search.best_estimator_.get_params()
        for param_dict in params[model_name]:
            for param_name, param_value in param_dict.items():
                print("\t%s: %r" % (param_name,
best_parameters[param_name]))

```

grid search fait
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.847

meilleur estimateur SVC(kernel='linear', random_state=42)

Accuracy : 0.870

Classification Report

	precision	recall	f1-score	support
FALSE	0.94203	0.78313	0.85526	83
TRUE	0.82000	0.95349	0.88172	86
accuracy			0.86982	169
macro avg	0.88101	0.86831	0.86849	169
weighted avg	0.87993	0.86982	0.86873	169

Ensemble des meilleurs paramètres :

```

C: 1.0
gamma: 'scale'
kernel: 'linear'

```

grid search fait
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.846

meilleur estimateur SVC(kernel='linear', random_state=42)

Accuracy : 0.846

Classification Report

	precision	recall	f1-score	support
FALSE	0.91304	0.75904	0.82895	83
TRUE	0.80000	0.93023	0.86022	86
accuracy			0.84615	169
macro avg	0.85652	0.84463	0.84458	169
weighted avg	0.85552	0.84615	0.84486	169

Ensemble des meilleurs paramètres :

```

C: 1.0
gamma: 'scale'
kernel: 'linear'

```

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.849

meilleur estimateur SVC(C=10, random_state=42)

Accuracy : 0.888

Classification Report

	precision	recall	f1-score	support
FALSE	0.91026	0.85542	0.88199	83
TRUE	0.86813	0.91860	0.89266	86
accuracy			0.88757	169
macro avg	0.88919	0.88701	0.88732	169
weighted avg	0.88882	0.88757	0.88742	169

Ensemble des meilleurs paramètres :

C: 10

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.849

meilleur estimateur SVC(C=10, random_state=42)

Accuracy : 0.893

Classification Report

	precision	recall	f1-score	support
FALSE	0.91139	0.86747	0.88889	83
TRUE	0.87778	0.91860	0.89773	86
accuracy			0.89349	169
macro avg	0.89459	0.89304	0.89331	169
weighted avg	0.89429	0.89349	0.89339	169

Ensemble des meilleurs paramètres :

C: 10

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.881

meilleur estimateur SVC(C=5, random_state=42)

Accuracy : 0.882

Classification Report

	precision	recall	f1-score	support
FALSE	0.87952	0.87952	0.87952	83
TRUE	0.88372	0.88372	0.88372	86
accuracy			0.88166	169
macro avg	0.88162	0.88162	0.88162	169
weighted avg	0.88166	0.88166	0.88166	169

Ensemble des meilleurs paramètres :

C: 5

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.884

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.899

Classification Report

	precision	recall	f1-score	support
FALSE	0.86667	0.93976	0.90173	83
TRUE	0.93671	0.86047	0.89697	86
accuracy			0.89941	169
macro avg	0.90169	0.90011	0.89935	169
weighted avg	0.90231	0.89941	0.89931	169

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.884

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.899

Classification Report

	precision	recall	f1-score	support
FALSE	0.88372	0.91566	0.89941	83
TRUE	0.91566	0.88372	0.89941	86
accuracy			0.89941	169
macro avg	0.89969	0.89969	0.89941	169

weighted avg 0.89998 0.89941 0.89941 169

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.881

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.888

Classification Report

	precision	recall	f1-score	support
FALSE	0.88095	0.89157	0.88623	83
TRUE	0.89412	0.88372	0.88889	86
accuracy			0.88757	169
macro avg	0.88754	0.88764	0.88756	169
weighted avg	0.88765	0.88757	0.88758	169

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.862

meilleur estimateur RandomForestClassifier(max_features='log2', random_state=42)

Accuracy : 0.870

Classification Report

	precision	recall	f1-score	support
FALSE	0.85882	0.87952	0.86905	83
TRUE	0.88095	0.86047	0.87059	86
accuracy			0.86982	169
macro avg	0.86989	0.86999	0.86982	169
weighted avg	0.87008	0.86982	0.86983	169

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.856

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.858

Classification Report

	precision	recall	f1-score	support
FALSE	0.89333	0.80723	0.84810	83
TRUE	0.82979	0.90698	0.86667	86
accuracy			0.85799	169
macro avg	0.86156	0.85710	0.85738	169
weighted avg	0.86100	0.85799	0.85755	169

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.865

meilleur estimateur RandomForestClassifier(max_features='log2',
random_state=42)

Accuracy : 0.876

Classification Report

	precision	recall	f1-score	support
FALSE	0.86047	0.89157	0.87574	83
TRUE	0.89157	0.86047	0.87574	86
accuracy			0.87574	169
macro avg	0.87602	0.87602	0.87574	169
weighted avg	0.87629	0.87574	0.87574	169

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.867

meilleur estimateur RandomForestClassifier(max_features='log2',
random_state=42)

Accuracy : 0.876

Classification Report

	precision	recall	f1-score	support
FALSE	0.82979	0.93976	0.88136	83
TRUE	0.93333	0.81395	0.86957	86
accuracy			0.87574	169
macro avg	0.88156	0.87686	0.87546	169
weighted avg	0.88248	0.87574	0.87536	169

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.862

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.882

Classification Report

	precision	recall	f1-score	support
FALSE	0.87952	0.87952	0.87952	83
TRUE	0.88372	0.88372	0.88372	86
accuracy			0.88166	169
macro avg	0.88162	0.88162	0.88162	169
weighted avg	0.88166	0.88166	0.88166	169

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.874

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.899

Classification Report

	precision	recall	f1-score	support
FALSE	0.90244	0.89157	0.89697	83
TRUE	0.89655	0.90698	0.90173	86
accuracy			0.89941	169

macro avg	0.89950	0.89927	0.89935	169
weighted avg	0.89944	0.89941	0.89939	169

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.873

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.911

Classification Report

	precision	recall	f1-score	support
FALSE	0.88636	0.93976	0.91228	83
TRUE	0.93827	0.88372	0.91018	86
accuracy			0.91124	169
macro avg	0.91232	0.91174	0.91123	169
weighted avg	0.91278	0.91124	0.91121	169

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.867

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.905



Classification Report



	precision	recall	f1-score	support
FALSE	0.90361	0.90361	0.90361	83
TRUE	0.90698	0.90698	0.90698	86
accuracy			0.90533	169
macro avg	0.90530	0.90530	0.90530	169
weighted avg	0.90533	0.90533	0.90533	169



Ensemble des meilleurs paramètres :



n_estimators: 50



max_features: 'sqrt'























On réapprends sur le X