

#CLASSIFICATION : TRUE/FALSE VS OTHER :

Membres: Hadjoudja Bachir (21811363), Zeggar Rym (21909615), Bendahmane Rania (21811387), Labiad Youcef (21710780).

#les imports utilisés dans ce notebook

```
import sys
from numpy import vstack
import pandas as pd
from pandas import read_csv
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from torch.utils.data import random_split
from torch import Tensor
from torch.nn import Linear
from torch.nn import ReLU
from torch.nn import Sigmoid
from torch.nn import Module
from torch.optim import SGD
from torch.nn import BCELoss
from torch.nn.init import kaiming_uniform_
from torch.nn.init import xavier_uniform_
import re
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from pandas import read_csv
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
import pickle
import string

import nltk
from nltk.stem import WordNetLemmatizer
from nltk.stem import PorterStemmer
from nltk.corpus import stopwords
from nltk import word_tokenize
from sklearn.pipeline import Pipeline
```

autorisation

```
from google.colab import drive
drive.mount('/content/gdrive/')
```

Drive already mounted at /content/gdrive/; to attempt to forcibly remount, call drive.mount("/content/gdrive/", force_remount=True).

chemin spécifique Google Drive

```

my_local_drive='/content/gdrive/My Drive/Colab Notebooks'
# Ajout du path pour les librairies, fonctions et données
sys.path.append(my_local_drive)
# Se positionner sur le répertoire associé
%cd $my_local_drive
%ls

%pwd

/content/gdrive/My Drive/Colab Notebooks
avecscaler.pkl
Classification_de_données_textuelles2023.ipynb
'Copie de TRUE_FALSE_vs_OTHER.ipynb'
Dataset/
firstmodel.pkl
'Ingénierie_des_données_textuelles2023 (1).ipynb'
Ingénierie_des_données_textuelles2023.ipynb
MyNLPUilities.py
newsTrain2.csv
newsTrain_-_newsTrain.csv
penguins.csv
penguins.csv.1
pkl_modelNB.sav
Premières_Classifications.ipynb
'Projet ML FakeNEWS_TRUE_FALSE_TEXT.ipynb'
'Projet ML FakeNEWS_TRUE_FALSE_TEXT+TITRE.ipynb'
'Projet ML FakeNEWS_TRUE_FALSE_TEXT+TITRE_TOPIC_MODELLING.ipynb'
'Projet ML FakeNEWS_TRUE_FALSE_TITRE.ipynb'
__pycache__/
ReviewsLabelled.csv
ReviewsLabelled.csv.1
ReviewsLabelled.csv.2
ReviewsLabelled.csv.3
ReviewsLabelled.csv.4
ReviewsLabelled.csv.5
SentimentModel.pkl
StopWordsFrench.csv
StopWordsFrench.csv.1
StopWordsFrench.csv.2
StopWordsFrench.csv.3
StopWordsFrench.csv.4
Topics_extraction.ipynb
TP1_HAI817I.ipynb
TP2_HAI817I.ipynb
'TRUE_FALSE_vs_OTHER.ipynb'
Visualisation_Donnees_2D_3D.ipynb

{"type": "string"}

```

La fonction qui sera utilisée pour les prétraitements: MyCleanText

- Mettre le texte en minuscule
- Se débarrasser des stopwords
- Se débarrasser des nombres
- Stemmatisation
- Lemmatisation ..

```
#.....Fonction
MyCleanText .....
.....
# mettre en minuscule
#enlever les stopwords
#se debarrasser des nombres
#stemmatisation
#lemmatisation
#.....
.....
.....
```

```
nltk.download('wordnet')
nltk.download('stopwords')
nltk.download('punkt')
#liste des stopwords en anglais
stop_words = set(stopwords.words('english'))
```

```
def MyCleanText(X,
                lowercase=False, #mettre en minuscule
                removestopwords=False, #supprimer les stopwords
                removedigit=False, #supprimer les nombres
                getstemmer=False, #conserver la racine des termes
                getlemmatisation=False #lemmatisation des termes
                ):
    #conversion du texte d'entrée en chaîne de caractères
    sentence=str(X)
    #suppression des caractères spéciaux
    sentence = re.sub(r'^\w\s',' ', sentence)
    # suppression de tous les caractères uniques
    sentence = re.sub(r'\s+[a-zA-Z]\s+', ' ', sentence)
    # substitution des espaces multiples par un seul espace
    sentence = re.sub(r'\s+', ' ', sentence, flags=re.I)

    # decoupage en mots
    tokens = word_tokenize(sentence)
    if lowercase:
        tokens = [token.lower() for token in tokens]

    # suppression ponctuation
    table = str.maketrans('', '', string.punctuation)
    words = [token.translate(table) for token in tokens]
```

```

# suppression des tokens non alphanumérique ou numérique
words = [word for word in words if word.isalnum()]

# suppression des tokens numérique
if removedigit:
    words = [word for word in words if not word.isdigit()]

# suppression des stopwords
if removestopwords:
    words = [word for word in words if not word in stop_words]

# lemmatisation
if getlemmatisation:
    lemmatizer=WordNetLemmatizer()
    words = [lemmatizer.lemmatize(word) for word in words]

# racinisation
if getstemmer:
    ps = PorterStemmer()
    words=[ps.stem(word) for word in words]

sentence= ' '.join(words)

return sentence

```

```

[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!

```

- La classe TextNormalizer qui contiendra la fonction MyCleanText.
- Fit_transform de mon corpus propre.

```

#.....Etape 1 :
prétraitement du
texte .....

#.....Class
TextNormalizer .....

#fit_transform de mon corpus propre
#.....
.....
.....

```

```

from sklearn.base import BaseEstimator, TransformerMixin

class TextNormalizer(BaseEstimator, TransformerMixin):
    def __init__(self,
                  removestopwords=False, # suppression des stopwords
                  lowercase=False, # passage en minuscule
                  removedigit=False, # supprimer les nombres
                  getstemmer=False, # racinisation des termes
                  getlemmatisation=False # lemmatisation des termes
                  ):

        self.lowercase=lowercase
        self.getstemmer=getstemmer
        self.removestopwords=removestopwords
        self.getlemmatisation=getlemmatisation
        self.removedigit=removedigit

    def transform(self, X, **transform_params):
        # Nettoyage du texte
        X=X.copy() # pour conserver le fichier d'origine
        return [MyCleanText(text,lowercase=self.lowercase,
                             getstemmer=self.getstemmer,
                             removestopwords=self.removestopwords,
                             getlemmatisation=self.getlemmatisation,
                             removedigit=self.removedigit) for text in
X]

    def fit(self, X, y=None, **fit_params):
        return self

    def fit_transform(self, X, y=None, **fit_params):
        return self.fit(X).transform(X)

    def get_params(self, deep=True):
        return {
            'lowercase':self.lowercase,
            'getstemmer':self.getstemmer,
            'removestopwords':self.removestopwords,
            'getlemmatisation':self.getlemmatisation,
            'removedigit':self.removedigit
        }

    def set_params (self, **parameters):
        for parameter, value in parameters.items():
            setattr(self,parameter,value)
        return self

```

##Etape 1 : Préparer les données

- Load et preparer les données à partir des 2 fichiers csv

- Sélectionner que les lignes où on a True, False ou Other
- Après en créant une nouvelle colonne "regrouped" si la valeur de la colonne rating est true ou bien false on mettra TRUE/FALSE sinon on laisse OTHER

#Ici je cherche à sélectionner que les labels TRUE/FALSE et OTHER, donc les LIGNES qui contiennent au rating TRUE,FALSE et OTHER uniquement, le reste on enlève

```
dftrain = pd.read_csv("/content/gdrive/MyDrive/Colab
Notebooks/newsTrain2.csv", names=['id','text','title','rating'],
header=0,sep=',', encoding='utf8')
dftrain.reset_index(drop = True, inplace = True)
```

```
dftrain2 = pd.read_csv("/content/gdrive/MyDrive/Colab
Notebooks/newsTrain_-_newsTrain.csv",
names=['id','text','title','rating'], header=0,sep=',',
encoding='utf8')
dftrain2.reset_index(drop = True, inplace = True)
```

concaténer les deux dataframes en ajoutant les lignes du deuxième à la fin du premier

```
dftrain = pd.concat([dftrain, dftrain2], ignore_index=True)
dftrain = dftrain.loc[dftrain['rating'].isin(['TRUE','FALSE',
'other'])]
```

#On crée une colonne regroupe qui va mettre dans les lignes là où a true ou bien false la valeur TRUE/FALSE et OTHER ça laisse

```
dftrain['regrouped'] = dftrain['rating'].apply(lambda x: 'TRUE/FALSE'
if x in ['TRUE', 'FALSE'] else 'OTHER')
```

#Quelques affichages pour aider à mieux visualiser nos données

```
print("Echantillon de mon dataset \n")
print(dftrain.sample(n=10))
print("\n")
print("Quelques informations importantes \n")
dftrain.info()
```

Echantillon de mon dataset

	id	text \
2374	d8c5eecd	Teachers who leave the profession for other jo...
986	d1741354	News "If you won't lead, th...
832	97b3e15c	Denying 2000 years of the Medieval Warm Period...
2425	7f8bf578	President Trump has sometimes claimed that sci...
1451	2963ac03	General Colin Powell's Chief of Staff, Col. Wi...
2356	e5a08193	Image copyright Getty Images Workers on zero-...
2064	2695016	He did none of those things. I've reviewed all...
2205	9f10a8a9	It was an accurate and judicious answer, so na...
993	88a75bcc	A 30-year old man from Kentucky underwent some...

1799 d2a52dd6

AMA Lied – How Many Died?

```

                                title rating
regrouped
2374 Relentless' workload forcing 'desperate' teach... other
OTHER
986 IT'S OFFICIAL: Brexit Britain WILL thrive out ... FALSE
TRUE/FALSE
832 Gov't Seeks to Control 'Disorderly' Internet P... FALSE
TRUE/FALSE
2425 Climate Change Is Complex. We've Got Answers t... TRUE
TRUE/FALSE
1451 General Colin Powell's Chief of Staff drops th... FALSE
TRUE/FALSE
2356 Workplace reforms 'will protect gig economy wo... FALSE
TRUE/FALSE
2064 Getting on with the job this week – Scottish N... TRUE
TRUE/FALSE
2205                                A 62% Top Tax Rate? other
OTHER
993 Months after being denied media credentials fo... FALSE
TRUE/FALSE
1799 Impeachment Lawyers opened Trump's second Impe... FALSE
TRUE/FALSE
```

Quelques informations importantes

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1812 entries, 0 to 2527
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id          1812 non-null   object
1   text        1812 non-null   object
2   title       1784 non-null   object
3   rating      1812 non-null   object
4   regrouped   1812 non-null   object
dtypes: object(5)
memory usage: 84.9+ KB
```

Le jeu de données étant déséquilibré, on a pensé à appliquer le downsampling pour équilibrer nos données. on sélectionne des lignes aléatoirement de TRUE/FALSE de telle sorte que le nombre de lignes de TRUE/FALSE soit = au nbr de lignes de Other. et on mélange le DataFrame.

#On applique du sous-échantillonnage (downsampling) : car on a plus de FALSE (578) que des TRUE (211)

Séparer les classes en deux dataframes

```

df_false_true = dftrain[dftrain['regrouped'] == 'TRUE/FALSE']
df_other = dftrain [dftrain['regrouped'] == 'OTHER']

# Sous-échantillonner la classe majoritaire (FALSE) pour obtenir un
# nombre égal d'échantillons pour chaque classe
df_subsampled = df_false_true.sample(n=len(df_other), random_state=42)

# Concaténer les deux dataframes
dftrain = pd.concat([df_subsampled, df_other])

# Mélanger aléatoirement les données
dftrain = dftrain.sample(frac=1, random_state=42)

X_text=dftrain.iloc[0:,1:2]
X_title=dftrain.iloc[0:,2:3]

print("le texte est")
display(X_text)
print("le titre est")
display(X_title)

y=dftrain.iloc[0:,-1]
print("le y est")
display(y)
print("la taille de X_text est",X_text.shape)
print("la taille de y_train est ",y.shape)
print("les valeurs de TRUE/FALSE et OTHER maintenant sont
",y.value_counts())

```

le texte est

```

text
947 War-torn eastern regions of Ukraine have no la...
2224 TIJUANA, Mexico – It's the image from the unfo...
1307 Today, Congresswoman Maxine Waters D-CA, Chair...
798 Meghan Markle will use the furore over her int...
320 Further proof that Democrats are the greatest ...
...
1160 The scale of Antarctica is startling. Miles of...
570 Coronavirus may be sexually transmitted and ca...
1200 Like what? Helen Harwatt is a researcher trai...
2190 Tumeric kills cancer not patient
391 WASHINGTON, DC – The Pentagon has issued an in...

```

[468 rows x 1 columns]

le titre est

```

title
947 Look No Further, The Best Doctor Strange in th...

```



```

2224 A discussion of 'smokers' black lungs' started...
1307 Democratic Lawmaker introduces bill to rename ...
798  Newton Emerson: Swiss model offers food for th...
320  Democrats Introduce Bill To 'Euthanize Seniors...
...
1160          Miles of Ice Collapsing Into the Sea
570  Universal Credit leaves working families worse...
1200          If Everyone Ate Beans Instead of Beef
2190 Vermont state trooper revived with Narcan afte...
391  Pentagon Confirms Coronavirus Accidentally Got I...

```

[468 rows x 1 columns]

le y est

```

947    TRUE/FALSE
2224    TRUE/FALSE
1307    TRUE/FALSE
798      OTHER
320    TRUE/FALSE

```

```

...
1160    TRUE/FALSE
570      OTHER
1200      OTHER
2190      OTHER
391    TRUE/FALSE

```

Name: regrouped, Length: 468, dtype: object

la taille de X_text est (468, 1)

la taille de y_train est (468,)

les valeurs de TRUE et FALSE maintenant sont TRUE/FALSE 234
OTHER 234

Name: regrouped, dtype: int64

On divise notre grand X en jeu de données d'apprentissage et de test (20% de test).

```

X=dftrain.iloc[0:, 1:4]
print(X)

```

```

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size = 0.2,
random_state=10)
print("X_train is",X_train)
print("y_train is",y_train)
print("X_test is",X_test)
print("y_test is",y_test)

```

```

text \
947  War-torn eastern regions of Ukraine have no la...
2224 TIJUANA, Mexico – It's the image from the unfo...
1307 Today, Congresswoman Maxine Waters D-CA, Chair...

```

798 Meghan Markle will use the furore over her int...
 320 Further proof that Democrats are the greatest ...
 ...
 1160 The scale of Antarctica is startling. Miles of...
 570 Coronavirus may be sexually transmitted and ca...
 1200 Like what? Helen Harwatt is a researcher trai...
 2190 Tumeric kills cancer not patient
 391 WASHINGTON, DC – The Pentagon has issued an in...

	title	rating
947	Look No Further, The Best Doctor Strange in th...	FALSE
2224	A discussion of 'smokers' black lungs' started...	TRUE
1307	Democratic Lawmaker introduces bill to rename ...	FALSE
798	Newton Emerson: Swiss model offers food for th...	other
320	Democrats Introduce Bill To 'Euthanize Seniors...	FALSE
...		...
1160	Miles of Ice Collapsing Into the Sea	TRUE
570	Universal Credit leaves working families worse...	other
1200	If Everyone Ate Beans Instead of Beef	other
2190	Vermont state trooper revived with Narcan afte...	other
391	Pentagon Confirms Coronavirus Accidentally Got I...	FALSE

[468 rows x 3 columns]

X_train is

text \

2006 Historians may look to 2015 as the year when s...
 1834 Coronavirus may be sexually transmitted and ca...
 530 Contractors bidding for work with the governme...
 564 More CO2 would actually help the planet , says...
 340 To say out-loud that you find the results of t...
 ...
 1775 Last week, in the days leading up to Sanders' ...
 208 This is one in a series of articles taken from...
 562 Food parcels arriving at the Community Service...
 676 On Tuesday, radio show host John Fredricks sta...
 1925 A South African pastor, Alfred Ndlovu has died...

	title	rating
2006	NaN	other
1834	Universal Credit leaves working families worse...	other
530	Firms bidding for government contracts asked i...	other
564	Mitt Romney transfers \$1 million left over fro...	other
340	Reasons why the 2020 presidential election is ...	other
...		...
1775	Bernie makes it official: It's Biden or bust	other
208	European royals killing naked children for fun...	other
562	Adam Castillejo is still free of the virus mor...	other
676	Warren Statement on Boeing	other
1925	Joy Covey: Amazon pioneer and high tech rock s...	other

[374 rows x 3 columns]

y_train is 2006 OTHER

1834 OTHER

530 OTHER

564 OTHER

340 OTHER

...

1775 OTHER

208 OTHER

562 OTHER

676 OTHER

1925 OTHER

Name: regrouped, Length: 374, dtype: object

X_test is

text \

459 Thank you to Universities UK UUK for hosting u...

1519 Can any government statistics on COVID-19 deat...

925 MOSCOW – Russian President Vladimir Putin over...

1857 Please enable cookies on your web browser in o...

2386 PUPILS aged just five have been accused of sex...

...

1778 With a smile on her face, City Clerk Susana Me...

2205 It was an accurate and judicious answer, so na...

1568 Barack Obama, a former President of the US, wa...

66 Pennsylvania rejects 372,000 mail-in ballots, ...

2355 Rises in National Insurance Contributions NICS...

		title	rating
459	USCIS Announces Final Rule Enforcing Long-Stan...		TRUE
1519	The CDC Confesses to Lying About COVID-19 Deat...		FALSE
925	Short breaks damage young people's futures		other
1857	Denying 2000 years of the Medieval Warm Period...		FALSE
2386	Pervs' aged five School sex crime claims trebl...		other
...			...
1778	Trump administration asks Supreme Court to str...		TRUE
2205	A 62% Top Tax Rate?		other
1568	Former President Barack Obama arrested for ESP...		FALSE
66	Pennsylvania rejects 372,000 mail-in ballots, ...		FALSE
2355	Budget 2017: National Insurance rate rise crit...		other

[94 rows x 3 columns]

y_test is 459 TRUE/FALSE

1519 TRUE/FALSE

925 OTHER

1857 TRUE/FALSE

2386 OTHER

...

1778 TRUE/FALSE

2205 OTHER

1568 TRUE/FALSE

66 TRUE/FALSE

2355 OTHER
Name: regrouped, Length: 94, dtype: object

##Etape 2 : Classification selon la colonne TEXT :

Tester avec plusieurs classifieurs classiques.

```
import sklearn
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_fscore_support as score
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
```

```
# Importation des différentes librairies utiles pour le notebook
#Sickit learn met régulièrement à jour des versions et
#indique des futurs warnings.
#ces deux lignes permettent de ne pas les afficher.
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import sys
import pandas as pd
import numpy as np
import sklearn
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_fscore_support as score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
#Sickit learn met régulièrement à jour des versions et indique des
futurs warnings.
#ces deux lignes permettent de ne pas les afficher.
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)

from sklearn.metrics._plot.confusion_matrix import
ConfusionMatrixDisplay
from sklearn import metrics
from sklearn.metrics import confusion_matrix , ConfusionMatrixDisplay
from sklearn.metrics import classification_report

from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
import time

```

La fonction MyshowAllScores prend le y_test et le y_predict, affiche l'accuracy et le classification report avec la matrice de confusion.

```

# fonction qui affiche le classification report et la matrice de
confusion
def MyshowAllScores(y_test,y_pred):
    classes= np.unique(y_test)
    print("Accuracy : %0.3f"%(accuracy_score(y_test,y_pred)))
    print("Classification Report")
    print(classification_report(y_test,y_pred,digits=5))
    cnf_matrix = confusion_matrix(y_test,y_pred)
    disp=ConfusionMatrixDisplay(cnf_matrix,display_labels=classes)
    disp.plot()

```

Vu qu'on va travailler sur la colonne text, on va sélectionner cette dernière depuis le X_train et X_test pour apprendre et tester après.

```
X_train_text=X_train['text']
X_train_text.reset_index(drop = True, inplace = True)
X_test_text=X_test['text']
X_test_text.reset_index(drop = True, inplace = True)
```

Ici, c'est une étape importante, on va tester différents classifieurs, pour chacun des classifieurs, on va appliquer le prétraitement + Vectorisation Tfidf, et on applique une cross_val_score avec un Kfold de 10 fois, par la suite on stocke dans une liste all_results la moyenne des accuracy + l'écart type et on la trie par ordre décroissant de moyenne d'accuracy et d'écart type. on remarque que les 3 meilleurs sont SVM,LR et RF qu'on va sélectionner pour leur appliquer le GridSearch sur les paramètres des prétraitements + leurs hyperparamètres pour pouvoir choisir le meilleur.

```
np.random.seed(42) # Set the random seed for NumPy
```

```
score = 'accuracy'
seed = 7
allresults = []
results = []
names = []
```

```
# Liste des modèles à tester
```

```
models = [
    ('MultinomialNB', MultinomialNB()),
    ('LogisticRegression', LogisticRegression(random_state=42)),
    ('KNN', KNeighborsClassifier()),
    ('CART', DecisionTreeClassifier(random_state=42)),
    ('RF', RandomForestClassifier(random_state=42)),
    ('SVM', SVC(random_state=42))
]
```

```
# Création d'un pipeline pour chaque modèle
```

```
pipelines = []
for name,model in models:
    pipeline = Pipeline([
        ('normalize', TextNormalizer()),
        ('tfidf', TfidfVectorizer()),
        (name,model)
    ])
    pipelines.append((name,pipeline))
```

```
all_results=[]
scores=[]
names=[]
for p in pipelines:
    print(p[1])
    # cross validation en 10 fois
```

```

kfold = KFold(n_splits=10, random_state=seed, shuffle=True)
start_time = time.time()
# application de la classification
cv_results = cross_val_score(p[1], X_train_text, y_train, cv=kfold,
scoring=score)
scores.append(cv_results)
names.append(p[0])
all_results.append((p[0], cv_results.mean(), cv_results.std()))
end_time = time.time()

print("all resultats", all_results)

all_results = sorted(all_results, key=lambda x: (-x[1], -x[2]))
print("all resultats", all_results)

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('MultinomialNB', MultinomialNB())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('LogisticRegression',
LogisticRegression(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('KNN', KNeighborsClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('CART', DecisionTreeClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('RF', RandomForestClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
                ('SVM', SVC(random_state=42))])
all resultats [('MultinomialNB', 0.7273115220483642,
0.08390022989848737), ('LogisticRegression', 0.8527738264580369,
0.05393319572545874), ('KNN', 0.620199146514936, 0.06317935636792352),
('CART', 0.7891891891891893, 0.06474231016745102), ('RF',
0.8475106685633002, 0.0793750031815335), ('SVM', 0.893172119487909,
0.05576555995138434)]
all resultats [('SVM', 0.893172119487909, 0.05576555995138434),
('LogisticRegression', 0.8527738264580369, 0.05393319572545874),
('RF', 0.8475106685633002, 0.0793750031815335), ('CART',
0.7891891891891893, 0.06474231016745102), ('MultinomialNB',
0.7273115220483642, 0.08390022989848737), ('KNN', 0.620199146514936,
0.06317935636792352)]

```


On affiche les accuracy de chaque classifieur, on remarque la médiane (en rouge) de chaque et l'écart type aussi.

```

import matplotlib.pyplot as plt
fig = plt.figure()
fig.suptitle('Comparaison des algorithmes')
ax = fig.add_subplot(111)
plt.boxplot(scores)
ax.set_xticklabels(names)

[Text(1, 0, 'MultinomialNB'),
 Text(2, 0, 'LogisticRegression'),
 Text(3, 0, 'KNN'),
 Text(4, 0, 'CART'),
 Text(5, 0, 'RF'),
 Text(6, 0, 'SVM')]

```



Choisir les meilleurs paramètres et hyperparamètres pour SVM, RF et LR :

On a un pipeline pour chaque prétraitement différent, on essaye pas mal (miniscule, lemmatisation, miniscule + lemmatisation..) et on stocke le fit_transform de nos X_train, X_test sur les pipelines dans des listes qui vont contenir tous les fit_transform des pipelines pour chaque classifieur, par la suite on parcourt ces listes là, on itère dessus, et chaque élément de la liste (train) va passer par le GridSearch et puis on prédit sur son correspondant dans la liste (test).


```

from sklearn.model_selection import GridSearchCV
from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from tabulate import tabulate

```

```

np.random.seed(42) # Set the random seed for NumPy

```

```

# le plus simple est de faire un test sur différents pipelines.
# pipeline de l'utilisation de CountVectorizer sur le texte avec
différents pré-traitements

```

```

CV_brut = Pipeline([('cleaner', TextNormalizer()),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])
CV_lowcase = Pipeline([('cleaner',
                        TextNormalizer(removestopwords=False, lowercase=True,

```

```

                        getstemmer=False, removedigit=False)),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])
CV_lowStop = Pipeline([('cleaner',
                        TextNormalizer(removestopwords=True, lowercase=True,

```

```

                        getstemmer=False, removedigit=False)),
                    ('count_vectorizer',
                     CountVectorizer(lowercase=False))])

```

```

CV_lowStopstem = Pipeline([('cleaner',
                            TextNormalizer(removestopwords=True, lowercase=True,

```

```

                            getstemmer=True, removedigit=False)),
                            ('count_vectorizer',
                             CountVectorizer(lowercase=False))])

```

```

# pipeline de l'utilisation de TfidfVectorizer avec différents pré-
traitements

```

```

TFIDF_brut = Pipeline([('cleaner', TextNormalizer()),
                       ('tfidf_vectorizer',
                        TfidfVectorizer(lowercase=False))])

```

```

TFIDF_lowcase = Pipeline([('cleaner',
                           TextNormalizer(removestopwords=False, lowercase=True,

```

```

getstemmer=False,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])
TFIDF_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

```

```

# Liste de tous les modeles à tester
all_models = [
    ("CV_brut", CV_brut),
    ("CV_lowercase", CV_lowercase),
    ("CV_lowStop", CV_lowStop),
    ("CV_lowStopstem",CV_lowStopstem),
    ("TFIDF_lowercase", TFIDF_lowercase),
    ("TFIDF_lowStop", TFIDF_lowStop),
    ("TFIDF_lowStopstem",TFIDF_lowStopstem),
    ("TFIDF_brut", TFIDF_brut)
]

```

```

#on crée des listes qui vont contenir les fit_transform des pipelines
sur le X_train et X_test
X_train_text_SVC = []
X_test_text_SVC = []

```

```

X_train_text_RandomForestClassifier = []
X_test_text_RandomForestClassifier = []

```

```

X_train_text_LogisticRegression = []
X_test_text_LogisticRegression = []

```

```

for name, pipeline in all_models :

```

```

X_train_text_SVC.append(pipeline.fit_transform(X_train_text).toarray()
)
    X_test_text_SVC.append(pipeline.transform(X_test_text).toarray())

```

```

X_train_text_RandomForestClassifier.append(pipeline.fit_transform(X_tr

```

```

ain_text).toarray())

X_test_text_RandomForestClassifier.append(pipeline.transform(X_test_text).toarray())

X_train_text_LogisticRegression.append(pipeline.fit_transform(X_train_text).toarray())

X_test_text_LogisticRegression.append(pipeline.transform(X_test_text).toarray())

models = {
    'SVC': SVC(random_state=42),
    'LogisticRegression': LogisticRegression(random_state=42),
    'RandomForestClassifier': RandomForestClassifier(random_state=42)
}

params = {'SVC': [{ 'C': [0.001, 0.01, 0.1, 1, 2, 5, 7, 10]},
                  { 'gamma': [0.001, 0.01, 0.1, 0.2, 0.3, 0.5, 0.7, 1]}],
          'RandomForestClassifier': [{ 'n_estimators': [10, 50, 100, 200, 300]},
                                     { 'max_features': ['auto', 'sqrt', 'log2']}],
          'LogisticRegression': [{ 'penalty': ['l1', 'l2', 'elasticnet', 'none']},
                                 { 'C': [0.001, 0.01, 0.1, 1, 10, 100]},
                                 { 'fit_intercept': [True, False]},
                                 { 'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']}],
          { 'max_iter': [100, 1000, 10000]}]
}

#On itère sur le dictionnaire des modèles
for model_name, model in models.items():
    score='accuracy'
    X_train_text = eval('X_train_text_' + model_name) #On crée une variable X_train_text qui est dynamique qui est la liste du fit_transform du X_train d'un classifieur
    X_test_text = eval('X_test_text_' + model_name) #On crée une variable X_test_text qui est dynamique qui est la liste du fit_transform du X_test d'un classifieur
    for i in range(len(X_train_text)): #on itère sur cette liste
        grid_search = GridSearchCV(model, params[model_name], n_jobs=-1, verbose=1, scoring=score) #on applique le GridSearch
        print("grid search fait")
        grid_search.fit(X_train_text[i], y_train)
        print ('meilleur score %.3f'%(grid_search.best_score_), '\n')
        print ('meilleur estimateur', grid_search.best_estimator_, '\n')

```

```

y_pred = grid_search.predict(X_test_text[i]) #on predict
MyshowAllScores(y_test,y_pred) #matrice de confusion report
classification

```

```

print("Ensemble des meilleurs paramètres :")
best_parameters = grid_search.best_estimator_.get_params()
for param_dict in params[model_name]:
    for param_name, param_value in param_dict.items():
        print("\t%s: %r" % (param_name,
best_parameters[param_name]))

```

grid search fait
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.901

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.5

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.901

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.5

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.5

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.5

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.880

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.882

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.880

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.880

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:

10 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_vali

```

ation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

-----
5 fits failed with the following error:
Traceback (most recent call last):
File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[
    nan 0.82623423      nan 0.7914955  0.7514955  0.80486486
 0.82086486 0.82623423 0.8129009  0.8049009  0.82623423 0.83686486
 0.82623423 0.82623423 0.82353153 0.78079279 0.7621982  0.82623423
 0.82623423 0.82623423]
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-


```
regression
  n_iter_i = _check_optimize_result(
```

meilleur score 0.837

meilleur estimateur LogisticRegression(fit_intercept=False,
random_state=42)

Accuracy : 0.851

Classification Report

	precision	recall	f1-score	support
OTHER	0.76471	0.95122	0.84783	41
TRUE/FALSE	0.95349	0.77358	0.85417	53
accuracy			0.85106	94
macro avg	0.85910	0.86240	0.85100	94
weighted avg	0.87115	0.85106	0.85140	94

Ensemble des meilleurs paramètres :

```
penalty: 'l2'
C: 1.0
fit_intercept: False
solver: 'lbfgs'
max_iter: 100
```

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
10 fits failed out of a total of 100.
```

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
```

```
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
-----
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.
```

```
    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[      nan  0.80738739      nan  0.81834234  0.74882883  0.80479279
  0.81279279  0.80738739  0.79142342  0.79409009  0.80738739  0.8181982
  0.80738739  0.80738739  0.80209009  0.7649009   0.73016216  0.80738739
  0.80738739  0.80738739]
    warnings.warn(
```

meilleur score 0.818

meilleur estimateur LogisticRegression(penalty='none', random_state=42)

Accuracy : 0.862

Classification Report

	precision	recall	f1-score	support
OTHER	0.78000	0.95122	0.85714	41
TRUE/FALSE	0.95455	0.79245	0.86598	53
accuracy			0.86170	94
macro avg	0.86727	0.87184	0.86156	94
weighted avg	0.87841	0.86170	0.86213	94

Ensemble des meilleurs paramètres :

```
penalty: 'none'  
C: 1.0  
fit_intercept: True  
solver: 'lbfgs'  
max_iter: 100
```

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:

10 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit

solver = _check_solver(self.solver, self.penalty, self.dual)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver

raise ValueError(

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit

solver = _check_solver(self.solver, self.penalty, self.dual)

File

```
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[          nan  0.81816216          nan  0.8074955  0.75135135  0.81545946
  0.8181982  0.81816216  0.8154955  0.81545946  0.81816216  0.81812613
  0.81816216  0.81816216  0.81823423  0.8181982  0.8074955  0.81816216
  0.81816216  0.81816216]
warnings.warn(
```

meilleur score 0.818

meilleur estimateur LogisticRegression(random_state=42,
solver='liblinear')

Accuracy : 0.851

Classification Report

	precision	recall	f1-score	support
OTHER	0.76471	0.95122	0.84783	41
TRUE/FALSE	0.95349	0.77358	0.85417	53
accuracy			0.85106	94
macro avg	0.85910	0.86240	0.85100	94
weighted avg	0.87115	0.85106	0.85140	94

Ensemble des meilleurs paramètres :

```
penalty: 'l2'
C: 1.0
fit_intercept: True
solver: 'liblinear'
max_iter: 100
```

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
10 fits failed out of a total of 100.
```

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

```

-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[      nan 0.81005405      nan 0.81272072 0.77005405 0.80482883
 0.82075676 0.81005405 0.80479279 0.79679279 0.81005405 0.8396036
 0.81005405 0.81005405 0.81279279 0.82612613 0.81275676 0.81005405
 0.81005405 0.81005405]
    warnings.warn(

```

meilleur score 0.840

meilleur estimateur LogisticRegression(fit_intercept=False,
random_state=42)

Accuracy : 0.819

Classification Report

	precision	recall	f1-score	support
OTHER	0.72222	0.95122	0.82105	41
TRUE/FALSE	0.95000	0.71698	0.81720	53
accuracy			0.81915	94
macro avg	0.83611	0.83410	0.81913	94
weighted avg	0.85065	0.81915	0.81888	94

Ensemble des meilleurs paramètres :

penalty: 'l2'
C: 1.0
fit_intercept: False
solver: 'lbfgs'
max_iter: 100

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:

10 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit

solver = _check_solver(self.solver, self.penalty, self.dual)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver

raise ValueError(

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[      nan  0.8289009      nan  0.84493694  0.51603604  0.51603604
  0.58821622  0.8289009  0.8396036  0.84493694  0.8289009  0.8261982
  0.8289009  0.8289009  0.8289009  0.8289009  0.8289009  0.8289009
  0.8289009  0.8289009 ]
    warnings.warn(

```

meilleur score 0.845

meilleur estimateur LogisticRegression(C=100, random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

```

penalty: 'l2'
C: 100
fit_intercept: True
solver: 'lbfgs'
max_iter: 100

```

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
10 fits failed out of a total of 100.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.
```

Below are more details about the failures:

```
-----
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got
l1 penalty.
```

```
-----
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got
elasticnet penalty.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
```



```
[      nan 0.84227027      nan 0.85827027 0.51603604 0.51603604
 0.60691892 0.84227027 0.86627027 0.8636036 0.84227027 0.81012613
 0.84227027 0.84227027 0.83956757 0.84227027 0.84227027 0.84227027
 0.84227027 0.84227027]
warnings.warn(
```

meilleur score 0.866

meilleur estimateur LogisticRegression(C=10, random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

penalty: 'l2'

C: 10

fit_intercept: True

solver: 'lbfgs'

max_iter: 100

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:

10 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit

```

    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

-----
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[      nan  0.85027027      nan  0.86356757  0.51603604  0.51603604
  0.62032432  0.85027027  0.87700901  0.88234234  0.85027027  0.81538739
  0.85027027  0.85027027  0.85027027  0.85027027  0.85027027  0.85027027
  0.85027027  0.85027027]
    warnings.warn(

```

meilleur score 0.882

meilleur estimateur LogisticRegression(C=100, random_state=42)

Accuracy : 0.862

Classification Report

	precision	recall	f1-score	support
OTHER	0.78000	0.95122	0.85714	41
TRUE/FALSE	0.95455	0.79245	0.86598	53
accuracy			0.86170	94
macro avg	0.86727	0.87184	0.86156	94
weighted avg	0.87841	0.86170	0.86213	94

Ensemble des meilleurs paramètres :

```
penalty: 'l2'  
C: 100  
fit_intercept: True  
solver: 'lbfgs'  
max_iter: 100
```

grid search fait

Fitting 5 folds for each of 20 candidates, totalling 100 fits

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:

10 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit

solver = _check_solver(self.solver, self.penalty, self.dual)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver

raise ValueError(

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

5 fits failed with the following error:

Traceback (most recent call last):

File

"/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File

"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit

solver = _check_solver(self.solver, self.penalty, self.dual)

```
File
"/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got elasticnet penalty.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite:
[          nan 0.8341982          nan 0.86371171 0.51603604 0.51603604
 0.58014414 0.8341982  0.86097297 0.86097297 0.8341982  0.8341982
 0.8341982  0.8341982  0.8341982  0.8341982  0.8341982  0.8341982
 0.8341982  0.8341982 ]
warnings.warn(
```

meilleur score 0.864

meilleur estimateur LogisticRegression(penalty='none',
random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

```
penalty: 'none'
C: 1.0
fit_intercept: True
solver: 'lbfgs'
max_iter: 100
```

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.864

meilleur estimateur RandomForestClassifier(max_features='log2',
random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.864

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.926

Classification Report

	precision	recall	f1-score	support
OTHER	0.88636	0.95122	0.91765	41
TRUE/FALSE	0.96000	0.90566	0.93204	53
accuracy			0.92553	94
macro avg	0.92318	0.92844	0.92484	94
weighted avg	0.92788	0.92553	0.92576	94

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.824

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.936

Classification Report

	precision	recall	f1-score	support
OTHER	0.87234	1.00000	0.93182	41
TRUE/FALSE	1.00000	0.88679	0.94000	53
accuracy			0.93617	94
macro avg	0.93617	0.94340	0.93591	94
weighted avg	0.94432	0.93617	0.93643	94

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.829

meilleur estimateur RandomForestClassifier(max_features='log2',
random_state=42)

Accuracy : 0.915

Classification Report

	precision	recall	f1-score	support
OTHER	0.86667	0.95122	0.90698	41
TRUE/FALSE	0.95918	0.88679	0.92157	53
accuracy			0.91489	94
macro avg	0.91293	0.91901	0.91427	94
weighted avg	0.91883	0.91489	0.91520	94

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.864

meilleur estimateur RandomForestClassifier(max_features='log2',
random_state=42)

Accuracy : 0.957

Classification Report

	precision	recall	f1-score	support
OTHER	0.95122	0.95122	0.95122	41
TRUE/FALSE	0.96226	0.96226	0.96226	53
accuracy			0.95745	94
macro avg	0.95674	0.95674	0.95674	94
weighted avg	0.95745	0.95745	0.95745	94

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_plot/
confusion_matrix.py:136: RuntimeWarning: More than 20 figures have
been opened. Figures created through the pyplot interface
(`matplotlib.pyplot.figure`) are retained until explicitly closed and
may consume too much memory. (To control this warning, see the rcParam
`figure.max_open_warning`). Consider using
`matplotlib.pyplot.close()`.
    fig, ax = plt.subplots()
```

meilleur score 0.853

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.853

meilleur estimateur RandomForestClassifier(max_features='log2',
random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

n_estimators: 100

max_features: 'log2'

grid search fait
Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.853

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.947


Classification Report



	precision	recall	f1-score	support
OTHER	0.89130	1.00000	0.94253	41
TRUE/FALSE	1.00000	0.90566	0.95050	53
accuracy			0.94681	94
macro avg	0.94565	0.95283	0.94651	94
weighted avg	0.95259	0.94681	0.94702	94



Ensemble des meilleurs paramètres :



n_estimators: 200



max_features: 'sqrt'

























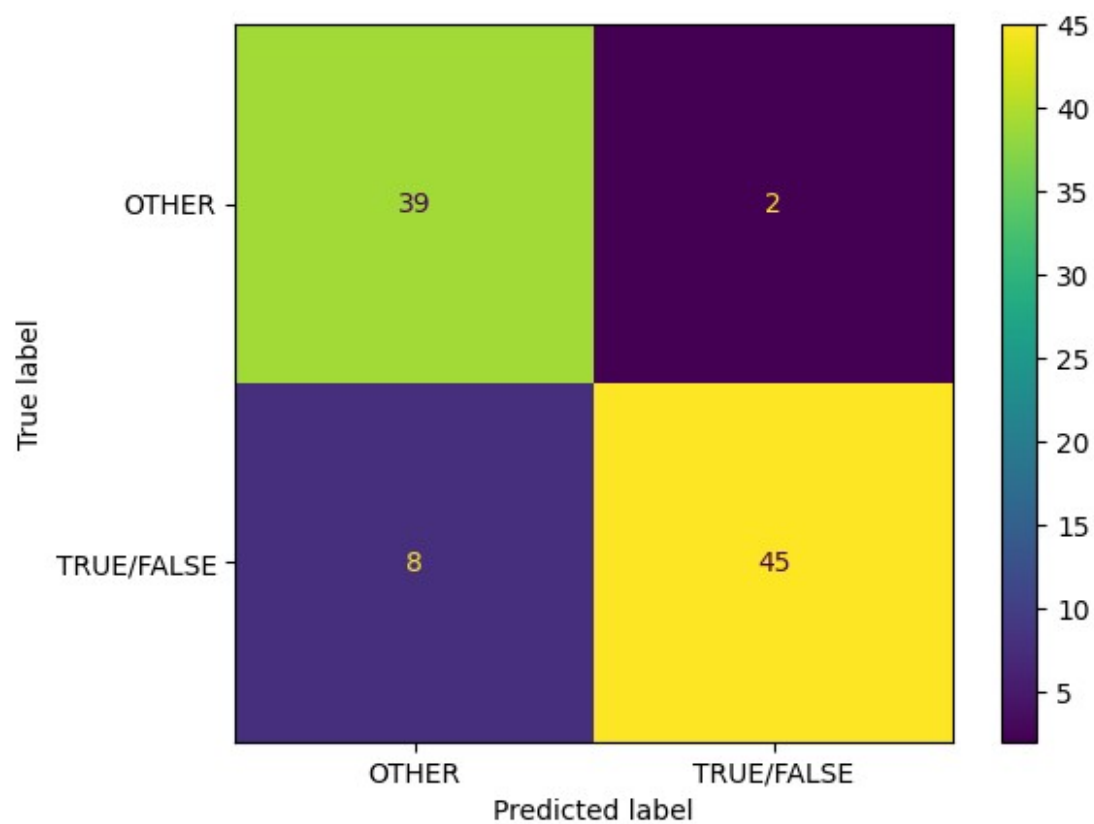
















##Etape 3 : Classification selon la colonne TITRE :

Vu qu'on va travailler sur la colonne titre, on va sélectionner cette dernière depuis le X_train et X_test pour apprendre et tester après.

```
X_train_title = X_train['title']
X_train_title.reset_index(drop = True, inplace = True)
X_test_title = X_test['title']
X_test_title.reset_index(drop = True, inplace = True)
```

Ici, c'est une étape importante, on va tester différents classifieurs, pour chacun des classifieurs, on va appliquer le prétraitement + Vectorisation Tfidf, et on applique une cross_val_score avec un Kfold de 10 fois, par la suite on stocke dans une liste all_results la moyenne des accuracy + l'écart type et on la trie par ordre décroissant de moyenne d'accuracy et d'écart type. on remarque que les 2 meilleurs sont SVM et RF qu'on va sélectionner pour leur appliquer le GridSearch sur les paramètres des prétraitements + leurs hyperparamètres pour pouvoir choisir le meilleur.

```
np.random.seed(42) # Set the random seed for NumPy
```

```
score = 'accuracy'
seed = 7
allresults = []
results = []
names = []
```

```

# Liste des modèles à tester
models = [
    ('MultinomialNB', MultinomialNB()),
    ('LogisticRegression', LogisticRegression(random_state=42)),
    ('KNN', KNeighborsClassifier()),
    ('CART', DecisionTreeClassifier(random_state=42)),
    ('RF', RandomForestClassifier(random_state=42)),
    ('SVM', SVC(random_state=42))
]

# Création d'un pipeline pour chaque modèle
pipelines = []
for name,model in models:
    pipeline = Pipeline([
        ('normalize', TextNormalizer()),
        ('tfidf', TfidfVectorizer()),
        (name,model)
    ])
    pipelines.append((name,pipeline))

all_results=[]
scores=[]
names=[]
for p in pipelines:
    print(p[1])
    # cross validation en 10 fois
    kfold = KFold(n_splits=10,random_state=seed,shuffle=True)
    start_time = time.time()
    # application de la classification
    cv_results = cross_val_score(p[1],X_train_title,y_train, cv=kfold,
scoring=score)
    scores.append(cv_results)
    names.append(p[0])
    all_results.append((p[0],cv_results.mean(),cv_results.std()))
    end_time = time.time()

all_results = sorted(all_results, key=lambda x: (-x[1], -x[2]))
print("all resultats", all_results)

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('MultinomialNB', MultinomialNB())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('LogisticRegression',
LogisticRegression(random_state=42))])

```

```

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('KNN', KNeighborsClassifier())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('CART', DecisionTreeClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('RF', RandomForestClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('SVM', SVC(random_state=42))])
all resultats [('SVM', 0.9172119487908962, 0.046935427835051634),
               ('RF', 0.82375533428165, 0.04554453420366589), ('LogisticRegression',
0.8211237553342817, 0.03992729512239782), ('MultinomialNB',
0.8130867709815078, 0.029869945063979022), ('CART',
0.8074679943100996, 0.061253246469652356), ('KNN', 0.6956614509246088,
0.07157923146711843)]

```

On affiche les accuracy de chaque classifieur, on remarque la médiane (en rouge) de chaque et l'écart type aussi.

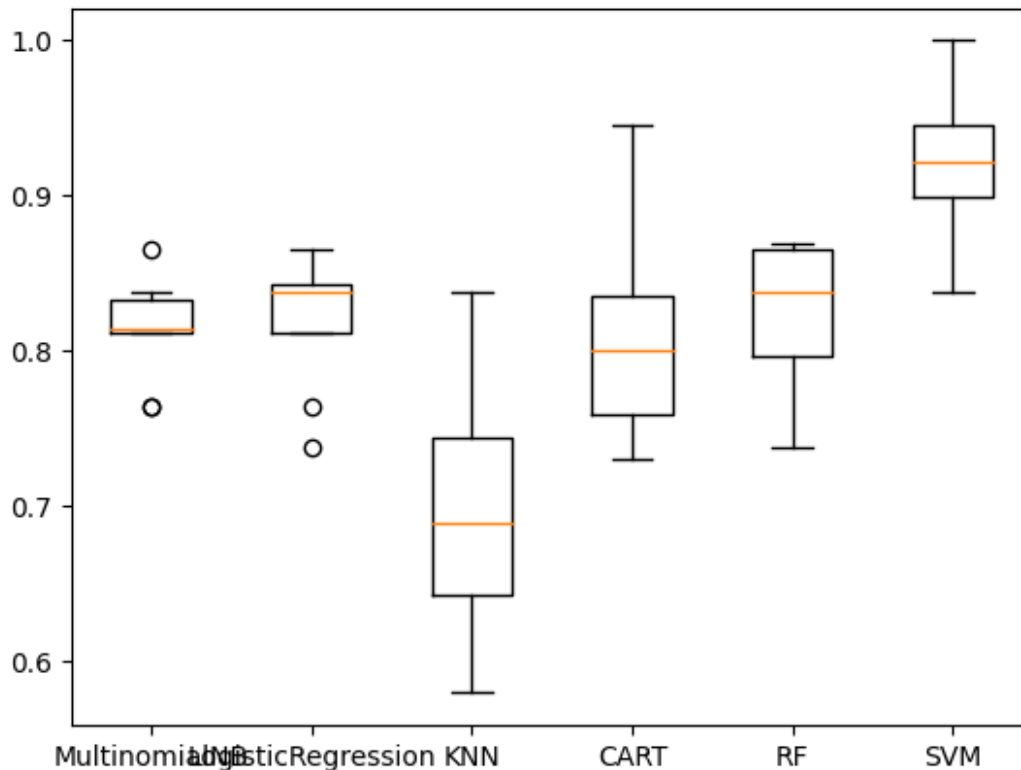
```

fig = plt.figure()
fig.suptitle('Comparaison des algorithmes')
ax = fig.add_subplot(111)
plt.boxplot(scores)
ax.set_xticklabels(names)

[Text(1, 0, 'MultinomialNB'),
 Text(2, 0, 'LogisticRegression'),
 Text(3, 0, 'KNN'),
 Text(4, 0, 'CART'),
 Text(5, 0, 'RF'),
 Text(6, 0, 'SVM')]

```

Comparaison des algorithmes



Choisir les meilleurs paramètres pour SVM et RF :

On a un pipeline pour chaque prétraitement différent, on essaye pas mal (miniscule, lemmatisation, miniscule + lemmatisation..) et on stocke le `fit_transform` de nos `X_train`, `X_test` sur les pipelines dans des listes qui vont contenir tous les `fit_transform` des pipelines pour chaque classifieur, par la suite on parcourt ces listes là, on itère dessus, et chaque élément de la liste (train) va passer par le `GridSearch` et puis on prédit sur son correspondant dans liste (test).

```
np.random.seed(42) # Set the random seed for NumPy
```

```
# le plus simple est de faire un test sur differents pipelines.
# pipeline de l'utilisation de CountVectorizer sur le texte avec
différents pre-traitements
```

[illegible]

```

CV_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
    ('count_vectorizer',
CountVectorizer(lowercase=False))])

CV_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
    ('count_vectorizer',
CountVectorizer(lowercase=False))])

# pipeline de l'utilisation de TfidfVectorizer avec differents pre-
traitements
TFIDF_brut = Pipeline ([('cleaner', TextNormalizer()),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowcase = Pipeline([('cleaner',
TextNormalizer(removestopwords=False,lowercase=True,

getstemmer=False,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])
TFIDF_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

# Liste de tous les modeles à tester
all_models = [
    ("CV_brut", CV_brut),
    ("CV_lowcase", CV_lowcase),
    ("CV_lowStop", CV_lowStop),
    ("CV_lowStopstem",CV_lowStopstem),
    ("TFIDF_lowcase", TFIDF_lowcase),
    ("TFIDF_lowStop", TFIDF_lowStop),

```



```

        ("TFIDF_lowStopstem",TFIDF_lowStopstem),
        ("TFIDF_brut", TFIDF_brut)
    ]

X_train_title_SVC = []
X_test_title_SVC = []

X_train_title_RandomForestClassifier = []
X_test_title_RandomForestClassifier = []

for name, pipeline in all_models :

X_train_title_SVC.append(pipeline.fit_transform(X_train_title).toarray
())

X_test_title_SVC.append(pipeline.transform(X_test_title).toarray())

X_train_title_RandomForestClassifier.append(pipeline.fit_transform(X_t
rain_title).toarray())

X_test_title_RandomForestClassifier.append(pipeline.transform(X_test_t
itle).toarray())

models = {
    'SVC': SVC(random_state=42),
    'RandomForestClassifier': RandomForestClassifier(random_state=42)
}

params = {'SVC': [{'C': [0.001, 0.01, 0.1, 1,2,5,7,10]},
                  {'gamma': [0.001, 0.01, 0.1,0.2,0.3,0.5,0.7,1]},
                  {'kernel': ['linear', 'rbf']}],
          'RandomForestClassifier': [{'n_estimators': [10, 50, 100, 200,
300]},
                                     {'max_features': ['auto', 'sqrt',
'log2']}]}

for model_name, model in models.items():
    score='accuracy'
    X_train_title = eval('X_train_title_' + model_name)
    X_test_title = eval('X_test_title_' + model_name)
    for i in range (len(X_train_title)):
        grid_search = GridSearchCV(model, params[model_name], n_jobs=-1,
verbose=1,scoring=score)
        print("grid search fait")
        grid_search.fit(X_train_title[i],y_train)

```

```

print ('meilleur score %0.3f'%(grid_search.best_score_),'\n')
print ('meilleur estimateur',grid_search.best_estimator_,'\n')
y_pred = grid_search.predict(X_test_title[i])
MyshowAllScores(y_test,y_pred)

print("Ensemble des meilleurs paramètres :")
best_parameters = grid_search.best_estimator_.get_params()
for param_dict in params[model_name]:
    for param_name, param_value in param_dict.items():
        print("\t%s: %r" % (param_name,
best_parameters[param_name]))

```

grid search fait
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.906

meilleur estimateur SVC(gamma=0.3, random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.3

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(gamma=0.3, random_state=42)

Accuracy : 0.904

Classification Report

	precision	recall	f1-score	support
OTHER	0.84783	0.95122	0.89655	41
TRUE/FALSE	0.95833	0.86792	0.91089	53
accuracy			0.90426	94
macro avg	0.90308	0.90957	0.90372	94
weighted avg	0.91013	0.90426	0.90464	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.3

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.898

meilleur estimateur SVC(gamma=0.7, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.7

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.898

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.936

Classification Report

	precision	recall	f1-score	support
OTHER	0.90698	0.95122	0.92857	41
TRUE/FALSE	0.96078	0.92453	0.94231	53
accuracy			0.93617	94
macro avg	0.93388	0.93787	0.93544	94
weighted avg	0.93732	0.93617	0.93632	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.5

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.904

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.957

Classification Report

	precision	recall	f1-score	support
OTHER	0.95122	0.95122	0.95122	41
TRUE/FALSE	0.96226	0.96226	0.96226	53
accuracy			0.95745	94
macro avg	0.95674	0.95674	0.95674	94
weighted avg	0.95745	0.95745	0.95745	94

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.898

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.904

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.797

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.830

Classification Report

	precision	recall	f1-score	support
OTHER	0.73585	0.95122	0.82979	41
TRUE/FALSE	0.95122	0.73585	0.82979	53
accuracy			0.82979	94
macro avg	0.84353	0.84353	0.82979	94

weighted avg 0.85728 0.82979 0.82979 94

Ensemble des meilleurs paramètres :

 n_estimators: 200

 max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.818

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.851

Classification Report

	precision	recall	f1-score	support
OTHER	0.74545	1.00000	0.85417	41
TRUE/FALSE	1.00000	0.73585	0.84783	53
accuracy			0.85106	94
macro avg	0.87273	0.86792	0.85100	94
weighted avg	0.88897	0.85106	0.85059	94

Ensemble des meilleurs paramètres :

 n_estimators: 50

 max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.842

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.883

Classification Report

	precision	recall	f1-score	support
OTHER	0.81250	0.95122	0.87640	41
TRUE/FALSE	0.95652	0.83019	0.88889	53
accuracy			0.88298	94
macro avg	0.88451	0.89070	0.88265	94
weighted avg	0.89370	0.88298	0.88344	94

Ensemble des meilleurs paramètres :

 n_estimators: 50

 max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.794

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.872

Classification Report

	precision	recall	f1-score	support
OTHER	0.79592	0.95122	0.86667	41
TRUE/FALSE	0.95556	0.81132	0.87755	53
accuracy			0.87234	94
macro avg	0.87574	0.88127	0.87211	94
weighted avg	0.88593	0.87234	0.87280	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.818

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.819

Classification Report

	precision	recall	f1-score	support
OTHER	0.70690	1.00000	0.82828	41
TRUE/FALSE	1.00000	0.67925	0.80899	53
accuracy			0.81915	94
macro avg	0.85345	0.83962	0.81864	94
weighted avg	0.87216	0.81915	0.81740	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.842

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.915

Classification Report

	precision	recall	f1-score	support
OTHER	0.86667	0.95122	0.90698	41
TRUE/FALSE	0.95918	0.88679	0.92157	53
accuracy			0.91489	94
macro avg	0.91293	0.91901	0.91427	94
weighted avg	0.91883	0.91489	0.91520	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.808

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.915

Classification Report

	precision	recall	f1-score	support
OTHER	0.86667	0.95122	0.90698	41
TRUE/FALSE	0.95918	0.88679	0.92157	53
accuracy			0.91489	94
macro avg	0.91293	0.91901	0.91427	94
weighted avg	0.91883	0.91489	0.91520	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.856

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.851

Classification Report


	precision	recall	f1-score	support
OTHER	0.74545	1.00000	0.85417	41
TRUE/FALSE	1.00000	0.73585	0.84783	53
accuracy			0.85106	94
macro avg	0.87273	0.86792	0.85100	94



weighted avg 0.88897 0.85106 0.85059 94



Ensemble des meilleurs paramètres :



n_estimators: 50



max_features: 'sqrt'


























##Etape 4 : Classification selon le TITRE ET TEXT ENSEMBLE (Concaténés):

- On va à partir de X_train concaténer les 2 colonnes TEXT et TITLE en mettant un espace entre les deux
- Vu qu'on va travailler sur la colonne text_titre qu'on vient de créer, on va sélectionner cette dernière depuis le X_train et X_test pour apprendre et tester après.

```
train_text_title = X_train.apply(lambda x : '{}
{}'.format(x['text'],x['title']),axis=1)
test_text_title = X_test.apply(lambda x : '{}
{}'.format(x['text'],x['title']),axis=1)
```

```
X_train['text_title'] = train_text_title
X_train_text_title = X_train['text_title']
X_train_text_title.reset_index(drop = True, inplace = True)
```

```
X_test['text_title'] = test_text_title
X_test_text_title = X_test['text_title']
X_test_text_title.reset_index(drop = True, inplace = True)
```

```
print("le texte et titre du train sont")
display(X_train_text_title)
```

```
print("le texte et titre du test sont")
display(X_test_text_title)
```

le texte et titre du train sont

```
0      Historians may look to 2015 as the year when s...
1      Coronavirus may be sexually transmitted and ca...
2      Contractors bidding for work with the governme...
3      More CO2 would actually help the planet , says...
4      To say out-loud that you find the results of t...
...
369    Last week, in the days leading up to Sanders' ...
370    This is one in a series of articles taken from...
371    Food parcels arriving at the Community Service...
372    On Tuesday, radio show host John Fredricks sta...
373    A South African pastor, Alfred Ndlovu has died...
Name: text_title, Length: 374, dtype: object
```

le texte et titre du test sont

```
0      Thank you to Universities UK UUK for hosting u...
1      Can any government statistics on COVID-19 deat...
2      MOSCOW – Russian President Vladimir Putin over...
3      Please enable cookies on your web browser in o...
4      PUPILS aged just five have been accused of sex...
...
89     With a smile on her face, City Clerk Susana Me...
90     It was an accurate and judicious answer, so na...
91     Barack Obama, a former President of the US, wa...
92     Pennsylvania rejects 372,000 mail-in ballots, ...
93     Rises in National Insurance Contributions NICS...
Name: text_title, Length: 94, dtype: object
```

Ici, c'est une étape importante, on va tester différents classifieurs, pour chacun des classifieurs, on va appliquer le prétraitement + Vectorisation Tfidf, et on applique une `cross_val_score` avec un Kfold de 10 fois, par la suite on stocke dans une liste `all_results` la moyenne des accuracy + l'écart type et on la trie par ordre décroissant de moyenne d'accuracy et d'écart type. on remarque que les 2 meilleurs sont SVM et RF qu'on va sélectionner pour leur appliquer le GridSearch sur les paramètres des prétraitements + leurs hyperparamètres pour pouvoir choisir le meilleur.

```
np.random.seed(42)  # Set the random seed for NumPy
```

```
score = 'accuracy'
seed = 7
allresults = []
results = []
names = []
```

```
# Liste des modèles à tester
```

```

models = [
    ('MultinomialNB', MultinomialNB()),
    ('LogisticRegression', LogisticRegression(random_state=42)),
    ('KNN', KNeighborsClassifier()),
    ('CART', DecisionTreeClassifier(random_state=42)),
    ('RF', RandomForestClassifier(random_state=42)),
    ('SVM', SVC(random_state=42))
]

# Création d'un pipeline pour chaque modèle
pipelines = []
for name,model in models:
    pipeline = Pipeline([
        ('normalize', TextNormalizer()),
        ('tfidf', TfidfVectorizer()),
        (name,model)
    ])
    pipelines.append((name,pipeline))

all_results=[]
scores=[]
names=[]
for p in pipelines:
    print(p[1])
    # cross validation en 10 fois
    kfold = KFold(n_splits=10,random_state=seed,shuffle=True)
    start_time = time.time()
    # application de la classification
    cv_results = cross_val_score(p[1],X_train_text_title,y_train,
cv=kfold, scoring=score)
    scores.append(cv_results)
    names.append(p[0])
    all_results.append((p[0],cv_results.mean(),cv_results.std()))
    end_time = time.time()

all_results = sorted(all_results, key=lambda x: (-x[1], -x[2]))
print("all resultats", all_results)

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('MultinomialNB', MultinomialNB())])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('LogisticRegression',
LogisticRegression(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('KNN', KNeighborsClassifier())])

```

```

Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('CART', DecisionTreeClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('RF', RandomForestClassifier(random_state=42))])
Pipeline(steps=[('normalize', TextNormalizer()), ('tfidf',
TfidfVectorizer()),
              ('SVM', SVC(random_state=42))])
all resultats [('SVM', 0.8876955903271693, 0.045964273722778934),
               ('RF', 0.8743954480796585, 0.05765956829705948),
               ('LogisticRegression', 0.8287339971550498, 0.0578545300951412),
               ('CART', 0.8155049786628734, 0.04056074962345757), ('MultinomialNB',
0.7514935988620198, 0.08135587683040285), ('KNN', 0.6389046941678521,
0.06555206867584254)]

```

On affiche les accuracy de chaque classifieur, on remarque la médiane (en rouge) de chaque et l'écart type aussi.


```

fig = plt.figure()
fig.suptitle('Comparaison des algorithmes')
ax = fig.add_subplot(111)
plt.boxplot(scores)
ax.set_xticklabels(names)

[Text(1, 0, 'MultinomialNB'),
 Text(2, 0, 'LogisticRegression'),
 Text(3, 0, 'KNN'),
 Text(4, 0, 'CART'),
 Text(5, 0, 'RF'),
 Text(6, 0, 'SVM')]

```

Comparaison des algorithmes



Choisir les meilleurs paramètres pour SVM et RF :

On a un pipeline pour chaque prétraitement différent, on essaye pas mal (miniscule, lemmatisation, miniscule + lemmatisation..) et on stocke le `fit_transform` de nos `X_train`, `X_test` sur les pipelines dans des listes qui vont contenir tous les `fit_transform` des pipelines pour chaque classifieur, par la suite on parcourt ces listes là, on itère dessus, et chaque élément de la liste (train) va passer par le `GridSearch` et puis on prédit sur son correspondant dans liste (test).

```
np.random.seed(42) # Set the random seed for NumPy
```

```
# le plus simple est de faire un test sur differents pipelines.
# pipeline de l'utilisation de CountVectorizer sur le texte avec
différents pre-traitements
```

[illegible]

```

CV_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
    ('count_vectorizer',
CountVectorizer(lowercase=False))])

CV_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
    ('count_vectorizer',
CountVectorizer(lowercase=False))])

# pipeline de l'utilisation de TfidfVectorizer avec differents pre-
traitements
TFIDF_brut = Pipeline ([('cleaner', TextNormalizer()),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowcase = Pipeline([('cleaner',
TextNormalizer(removestopwords=False,lowercase=True,

getstemmer=False,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])
TFIDF_lowStop = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=False,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

TFIDF_lowStopstem = Pipeline([('cleaner',
TextNormalizer(removestopwords=True,lowercase=True,

getstemmer=True,removedigit=False)),
    ('tfidf_vectorizer',
TfidfVectorizer(lowercase=False))])

# Liste de tous les modeles à tester
all_models = [
    ("CV_brut", CV_brut),
    ("CV_lowcase", CV_lowcase),
    ("CV_lowStop", CV_lowStop),
    ("CV_lowStopstem",CV_lowStopstem),
    ("TFIDF_lowcase", TFIDF_lowcase),
    ("TFIDF_lowStop", TFIDF_lowStop),

```

```

        ("TFIDF_lowStopstem",TFIDF_lowStopstem),
        ("TFIDF_brut", TFIDF_brut)
    ]

X_train_text_title_SVC = []
X_test_text_title_SVC = []

X_train_text_title_RandomForestClassifier = []
X_test_text_title_RandomForestClassifier = []

for name, pipeline in all_models :

X_train_text_title_SVC.append(pipeline.fit_transform(X_train_text_title).toarray())

X_test_text_title_SVC.append(pipeline.transform(X_test_text_title).toarray())

X_train_text_title_RandomForestClassifier.append(pipeline.fit_transform(X_train_text_title).toarray())

X_test_text_title_RandomForestClassifier.append(pipeline.transform(X_test_text_title).toarray())


models = {
    'SVC': SVC(random_state=42),
    'RandomForestClassifier': RandomForestClassifier(random_state=42)
}

params = {'SVC': [{'C': [0.001, 0.01, 0.1, 1,2,5,7,10]},
                  {'gamma': [0.001, 0.01, 0.1,0.2,0.3,0.5,0.7,1]},
                  {'kernel': ['linear', 'rbf']}]},
          'RandomForestClassifier': [{'n_estimators': [10, 50, 100, 200, 300]},
                                     {'max_features': ['auto', 'sqrt', 'log2']}]}

for model_name, model in models.items():
    score='accuracy'
    X_train_text_title = eval('X_train_text_title_' + model_name)
    X_test_text_title = eval('X_test_text_title_' + model_name)
    for i in range (len(X_train_text_title)):
        grid_search = GridSearchCV(model, params[model_name], n_jobs=-1,
        verbose=1,scoring=score)
        print("grid search fait")

```

```

grid_search.fit(X_train_text_title[i],y_train)
print ('meilleur score %0.3f'%(grid_search.best_score_),'\n')
print ('meilleur estimateur',grid_search.best_estimator_,'\n')
y_pred = grid_search.predict(X_test_text_title[i])
MyshowAllScores(y_test,y_pred)

print("Ensemble des meilleurs paramètres :")
best_parameters = grid_search.best_estimator_.get_params()
for param_dict in params[model_name]:
    for param_name, param_value in param_dict.items():
        print("\t%s: %r" % (param_name,
best_parameters[param_name]))

```

grid search fait
Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.901

meilleur estimateur SVC(gamma=0.2, random_state=42)

Accuracy : 0.979

Classification Report

	precision	recall	f1-score	support
OTHER	1.00000	0.95122	0.97500	41
TRUE/FALSE	0.96364	1.00000	0.98148	53
accuracy			0.97872	94
macro avg	0.98182	0.97561	0.97824	94
weighted avg	0.97950	0.97872	0.97865	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.2

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(gamma=0.2, random_state=42)

Accuracy : 0.979

Classification Report

	precision	recall	f1-score	support
OTHER	1.00000	0.95122	0.97500	41
TRUE/FALSE	0.96364	1.00000	0.98148	53
accuracy			0.97872	94
macro avg	0.98182	0.97561	0.97824	94

weighted avg 0.97950 0.97872 0.97865 94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.2

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(gamma=0.1, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.1

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits

meilleur score 0.901

meilleur estimateur SVC(gamma=0.1, random_state=42)

Accuracy : 0.968

Classification Report

	precision	recall	f1-score	support
OTHER	0.97500	0.95122	0.96296	41
TRUE/FALSE	0.96296	0.98113	0.97196	53
accuracy			0.96809	94
macro avg	0.96898	0.96618	0.96746	94
weighted avg	0.96821	0.96809	0.96804	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.1

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.874

meilleur estimateur SVC(C=2, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 2

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.869

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.869

meilleur estimateur SVC(gamma=0.5, random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
OTHER	0.84444	0.92683	0.88372	41
TRUE/FALSE	0.93878	0.86792	0.90196	53
accuracy			0.89362	94
macro avg	0.89161	0.89738	0.89284	94
weighted avg	0.89763	0.89362	0.89401	94

Ensemble des meilleurs paramètres :

C: 1.0

gamma: 0.5

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 18 candidates, totalling 90 fits
meilleur score 0.877

meilleur estimateur SVC(C=1, random_state=42)

Accuracy : 0.947

Classification Report

	precision	recall	f1-score	support
OTHER	0.92857	0.95122	0.93976	41
TRUE/FALSE	0.96154	0.94340	0.95238	53
accuracy			0.94681	94
macro avg	0.94505	0.94731	0.94607	94
weighted avg	0.94716	0.94681	0.94688	94

Ensemble des meilleurs paramètres :

C: 1

gamma: 'scale'

kernel: 'rbf'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits
meilleur score 0.840

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.936

Classification Report

	precision	recall	f1-score	support
OTHER	0.90698	0.95122	0.92857	41
TRUE/FALSE	0.96078	0.92453	0.94231	53
accuracy			0.93617	94
macro avg	0.93388	0.93787	0.93544	94

weighted avg 0.93732 0.93617 0.93632 94

Ensemble des meilleurs paramètres :

 n_estimators: 100

 max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.864

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.926

Classification Report

	precision	recall	f1-score	support
OTHER	0.88636	0.95122	0.91765	41
TRUE/FALSE	0.96000	0.90566	0.93204	53
accuracy			0.92553	94
macro avg	0.92318	0.92844	0.92484	94
weighted avg	0.92788	0.92553	0.92576	94

Ensemble des meilleurs paramètres :

 n_estimators: 200

 max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.815

meilleur estimateur RandomForestClassifier(random_state=42)

Accuracy : 0.904

Classification Report

	precision	recall	f1-score	support
OTHER	0.84783	0.95122	0.89655	41
TRUE/FALSE	0.95833	0.86792	0.91089	53
accuracy			0.90426	94
macro avg	0.90308	0.90957	0.90372	94
weighted avg	0.91013	0.90426	0.90464	94

Ensemble des meilleurs paramètres :

 n_estimators: 100

 max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.832

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.883

Classification Report

	precision	recall	f1-score	support
OTHER	0.81250	0.95122	0.87640	41
TRUE/FALSE	0.95652	0.83019	0.88889	53
accuracy			0.88298	94
macro avg	0.88451	0.89070	0.88265	94
weighted avg	0.89370	0.88298	0.88344	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.869

meilleur estimateur RandomForestClassifier(n_estimators=200,
random_state=42)

Accuracy : 0.926

Classification Report

	precision	recall	f1-score	support
OTHER	0.88636	0.95122	0.91765	41
TRUE/FALSE	0.96000	0.90566	0.93204	53
accuracy			0.92553	94
macro avg	0.92318	0.92844	0.92484	94
weighted avg	0.92788	0.92553	0.92576	94

Ensemble des meilleurs paramètres :

n_estimators: 200

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.856

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.894

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

OTHER	0.82979	0.95122	0.88636	41
TRUE/FALSE	0.95745	0.84906	0.90000	53
accuracy			0.89362	94
macro avg	0.89362	0.90014	0.89318	94
weighted avg	0.90177	0.89362	0.89405	94

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.837

meilleur estimateur RandomForestClassifier(n_estimators=300,
random_state=42)

Accuracy : 0.915

Classification Report

	precision	recall	f1-score	support
OTHER	0.86667	0.95122	0.90698	41
TRUE/FALSE	0.95918	0.88679	0.92157	53
accuracy			0.91489	94
macro avg	0.91293	0.91901	0.91427	94
weighted avg	0.91883	0.91489	0.91520	94

Ensemble des meilleurs paramètres :

n_estimators: 300

max_features: 'sqrt'

grid search fait

Fitting 5 folds for each of 8 candidates, totalling 40 fits

meilleur score 0.845

meilleur estimateur RandomForestClassifier(n_estimators=50,
random_state=42)

Accuracy : 0.883


Classification Report



	precision	recall	f1-score	support
OTHER	0.81250	0.95122	0.87640	41
TRUE/FALSE	0.95652	0.83019	0.88889	53
accuracy			0.88298	94
macro avg	0.88451	0.89070	0.88265	94
weighted avg	0.89370	0.88298	0.88344	94



Ensemble des meilleurs paramètres :



n_estimators: 50



max_features: 'sqrt'


























#Visualisation :

```
!pip install umap-learn[plot]
!pip install holoviews
!pip install -U ipykernel
```

```
from sklearn.decomposition import TruncatedSVD
import plotly.express as px
from sklearn.manifold import TSNE
# Umap
import umap.plot
from umap import UMAP
```

```
X_test_copy = X_test.copy()
```

```
tfidf=TfidfVectorizer()
vector_tfidf=tfidf.fit_transform(X_test_copy['text'])
```

2D

```
umap = UMAP(n_components=2, init='random', random_state=0)
projection = umap.fit_transform(vector_tfidf)
```

```
fig = px.scatter(
    projection, x=0, y=1,
    color=y_test, labels={'color': 'RATING'})
```

```
)
```

```
fig.show()
```